

15-721

DATABASE SYSTEMS



Lecture #06 – Indexing (Locking & Latching)

Andy Pavlo // Carnegie Mellon University // Spring 2016

TODAY'S AGENDA

Order Preserving Indexes

Index Locking & Latching

Prison Gang Tattoos

DATABASE INDEX

A data structure that improves the speed of data retrieval operations on a table at the cost of additional writes and storage space.

Indexes are used to quickly locate data without having to search every row in a table every time a table is accessed.

DATA STRUCTURES

Order Preserving Indexes

- A tree-like structure that maintains keys in some sorted order.
- Supports all possible predicates with $O(\log n)$ searches.

Hashing Indexes

- An associative array that maps a hash of the key to a particular record.
- Only supports equality predicates with $O(1)$ searches.

DATA STRUCTURES

Order Preserving Indexes

- A tree-like structure that maintains keys in some sorted order.
- Supports all possible predicates with $O(\log n)$ searches.

Hashing Indexes

- An associative array that maps a hash of the key to a particular record.
- Only supports equality predicates with $O(1)$ searches.

B-TREE VS. B+TREE


The original **B-tree** from 1972 stored keys + values in all nodes in the tree.

→ More memory efficient since each key only appears once in the tree.


A **B+tree** only stores values in leaf nodes. Inner nodes only guide the search process.

→ Easier to manage concurrent index access when the values are only in the leaf nodes.


B+TREE



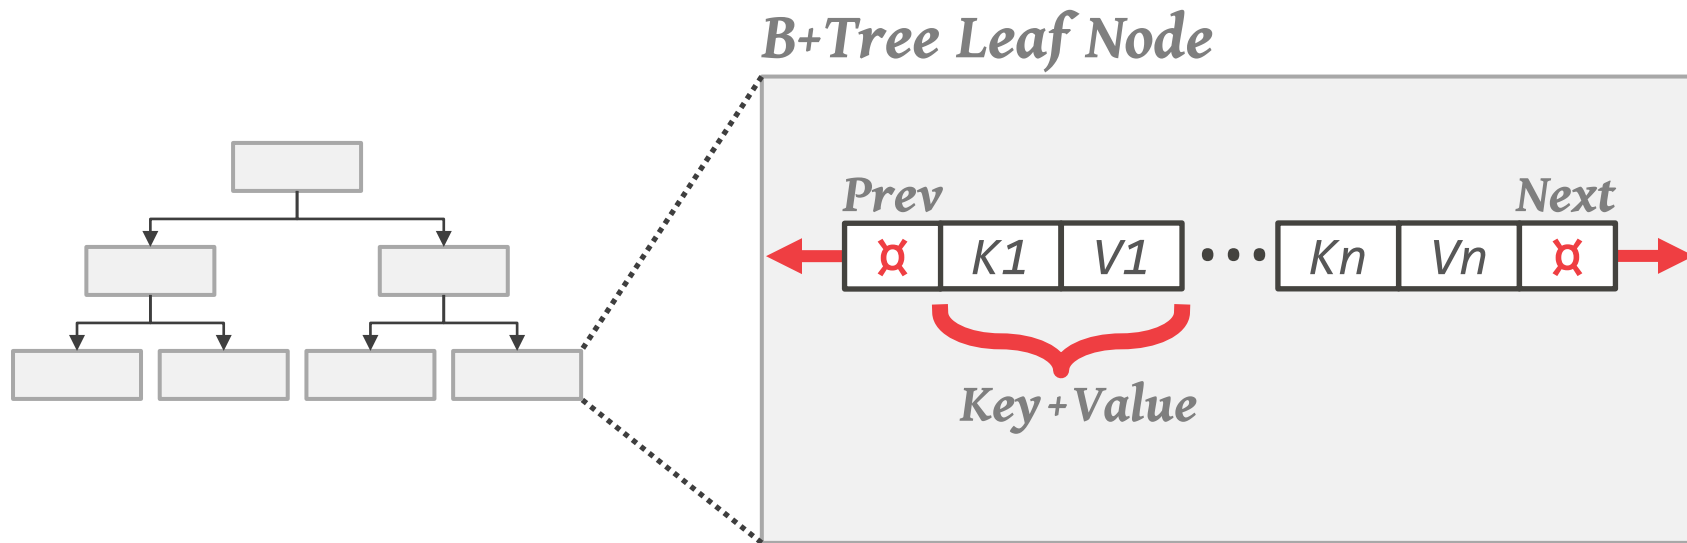
B+TREE



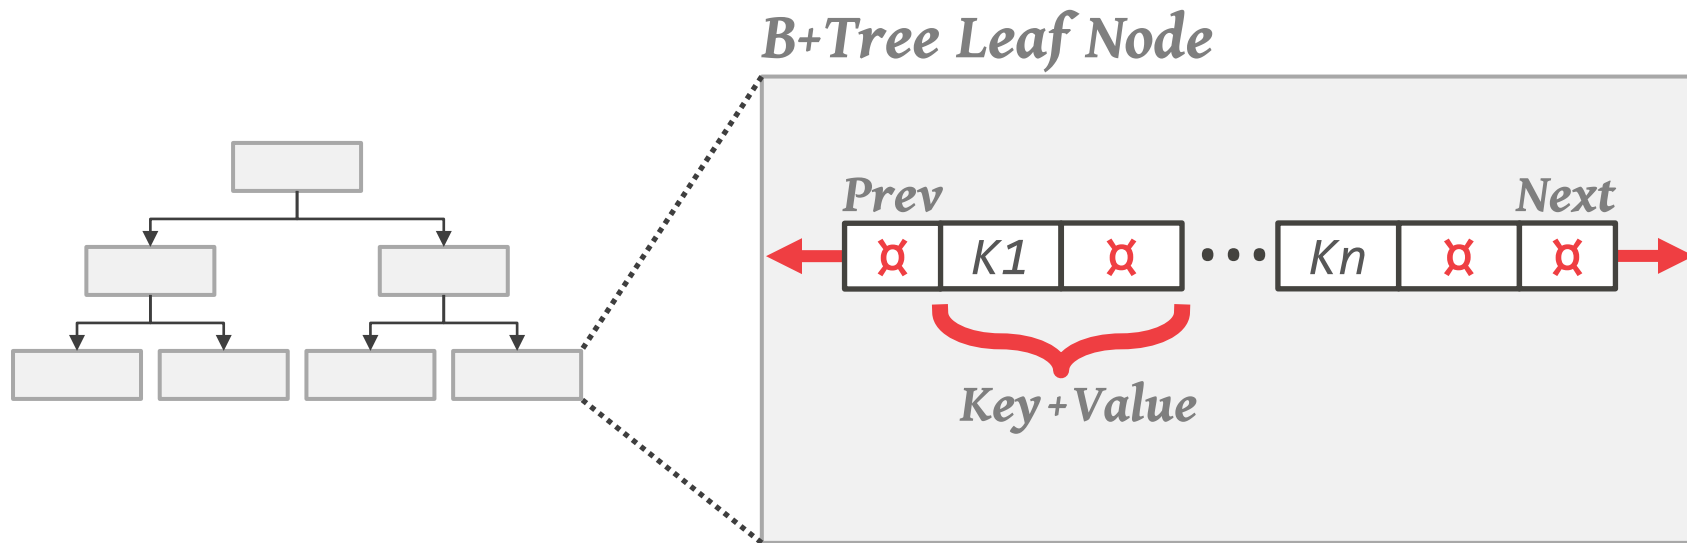
B+TREE



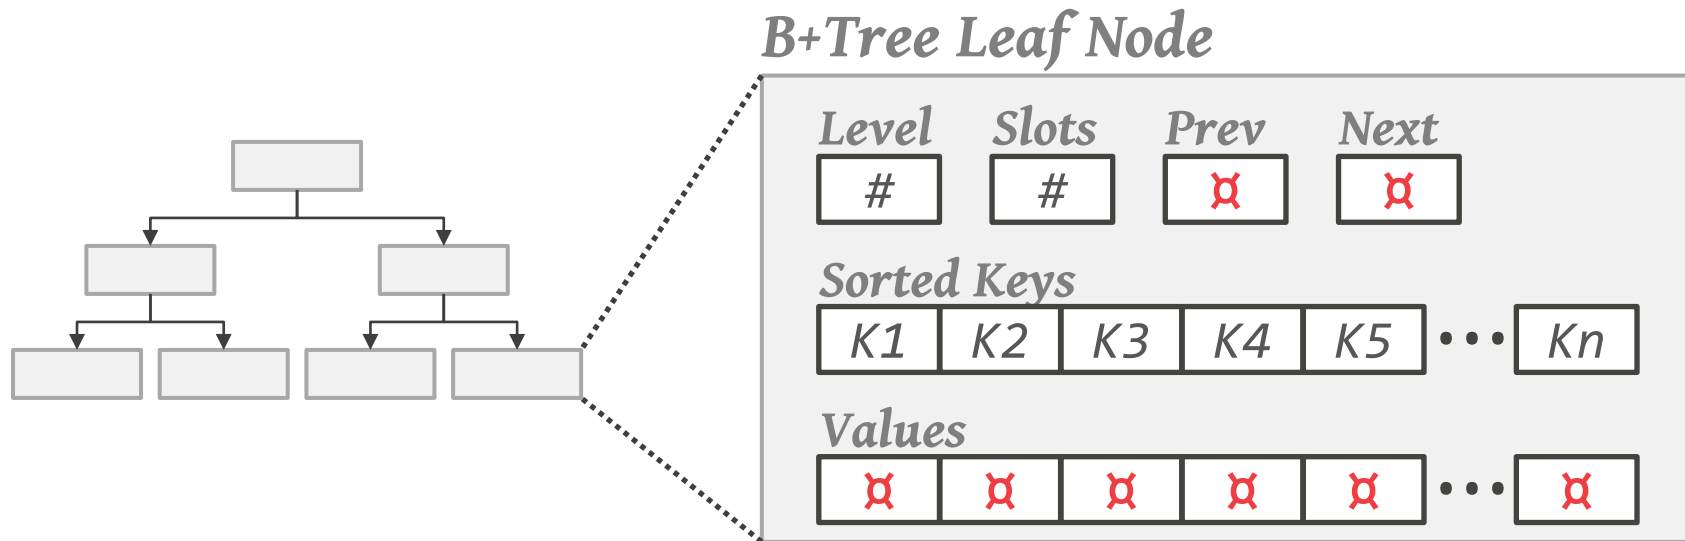
B+TREE




B+TREE



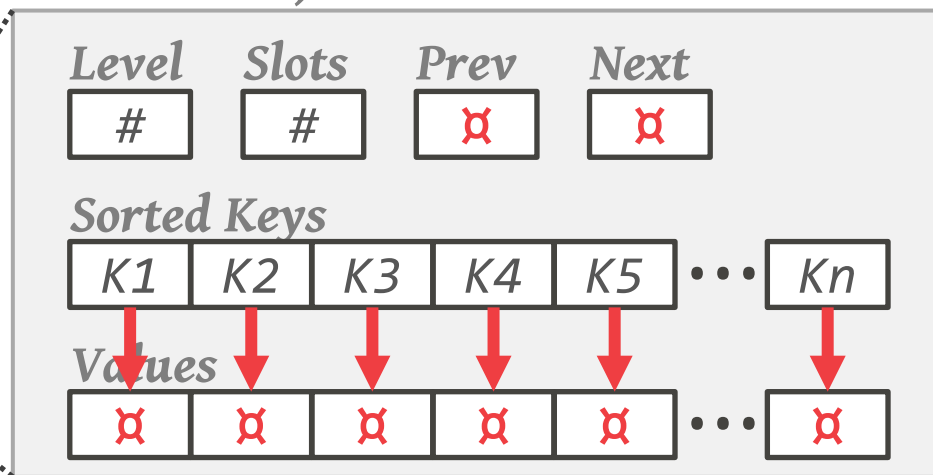
B+TREE



B+TREE



B+Tree Leaf Node



B+TREE DESIGN CHOICES

Non-Unique Indexes: One key maps to multiple values.

Variable Length Keys: The size of each key is not the same.

B+TREE: NON-UNIQUE INDEXES

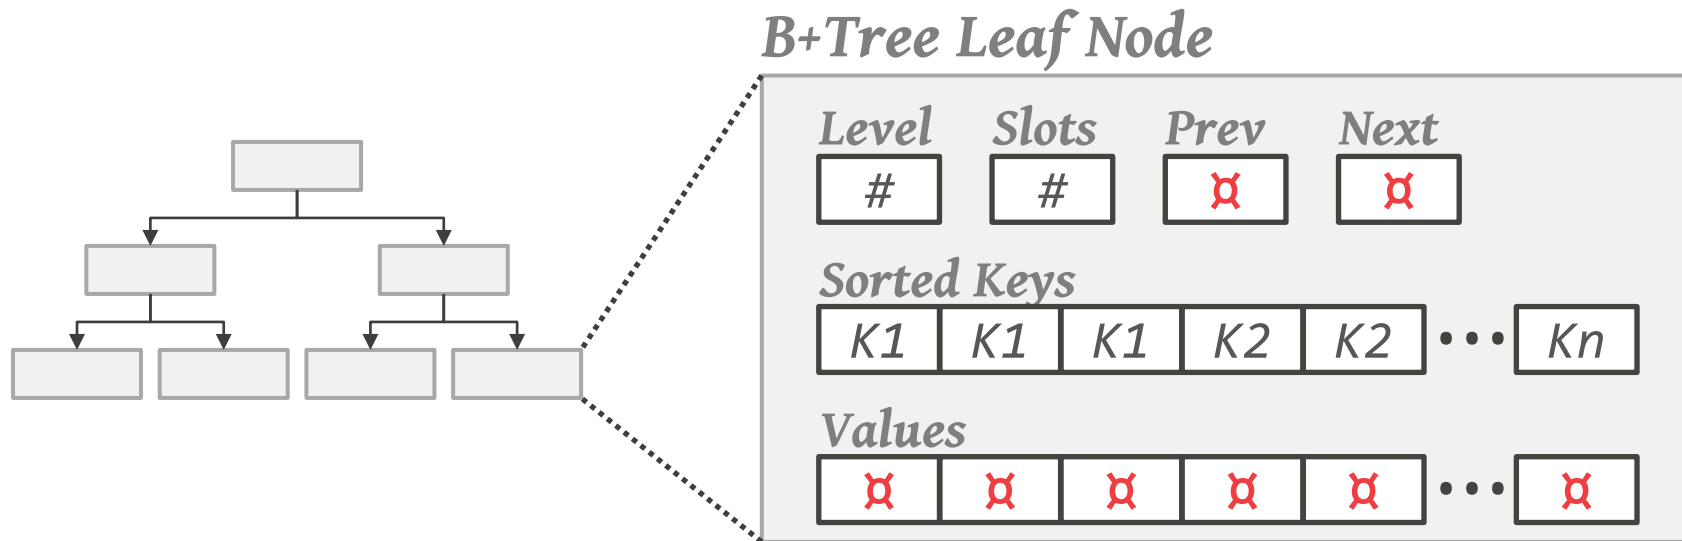
Approach #1: Duplicate Keys

- Use the same leaf node layout but store duplicate keys multiple times.

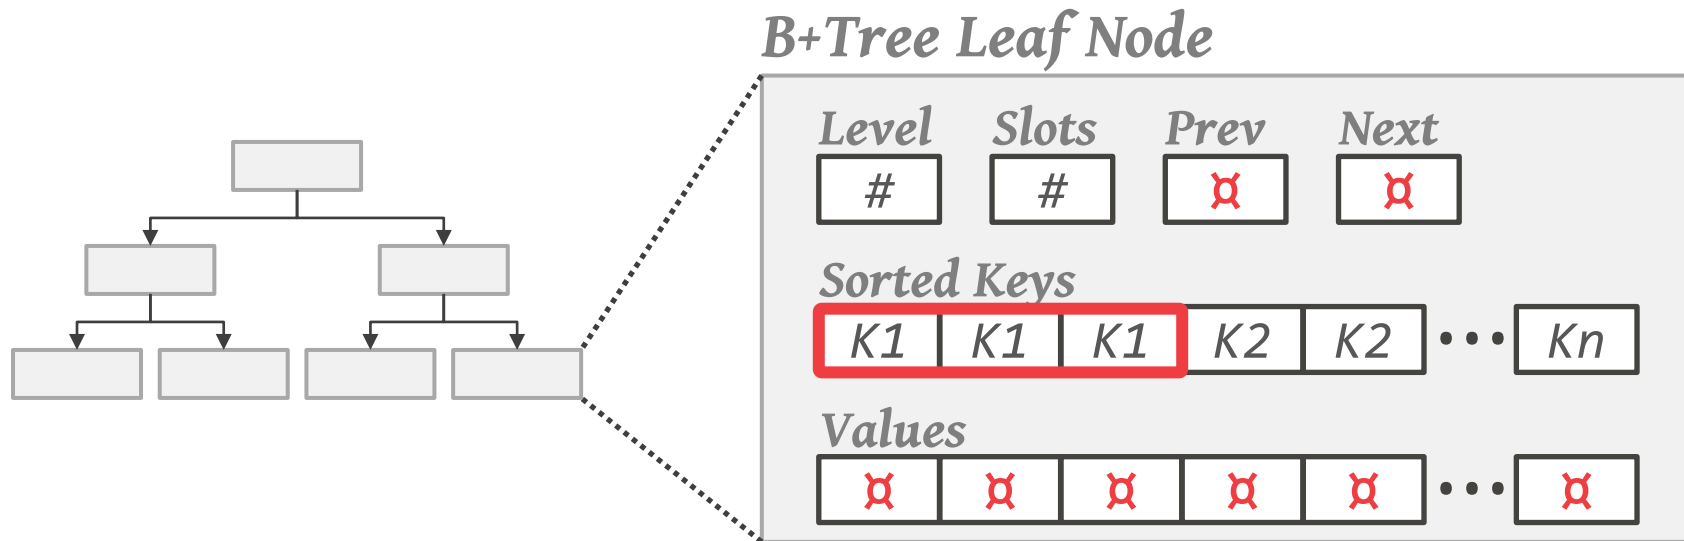
Approach #2: Value Lists

- Store each key only once and maintain a linked list of unique values.


B+TREE: DUPLICATE KEYS



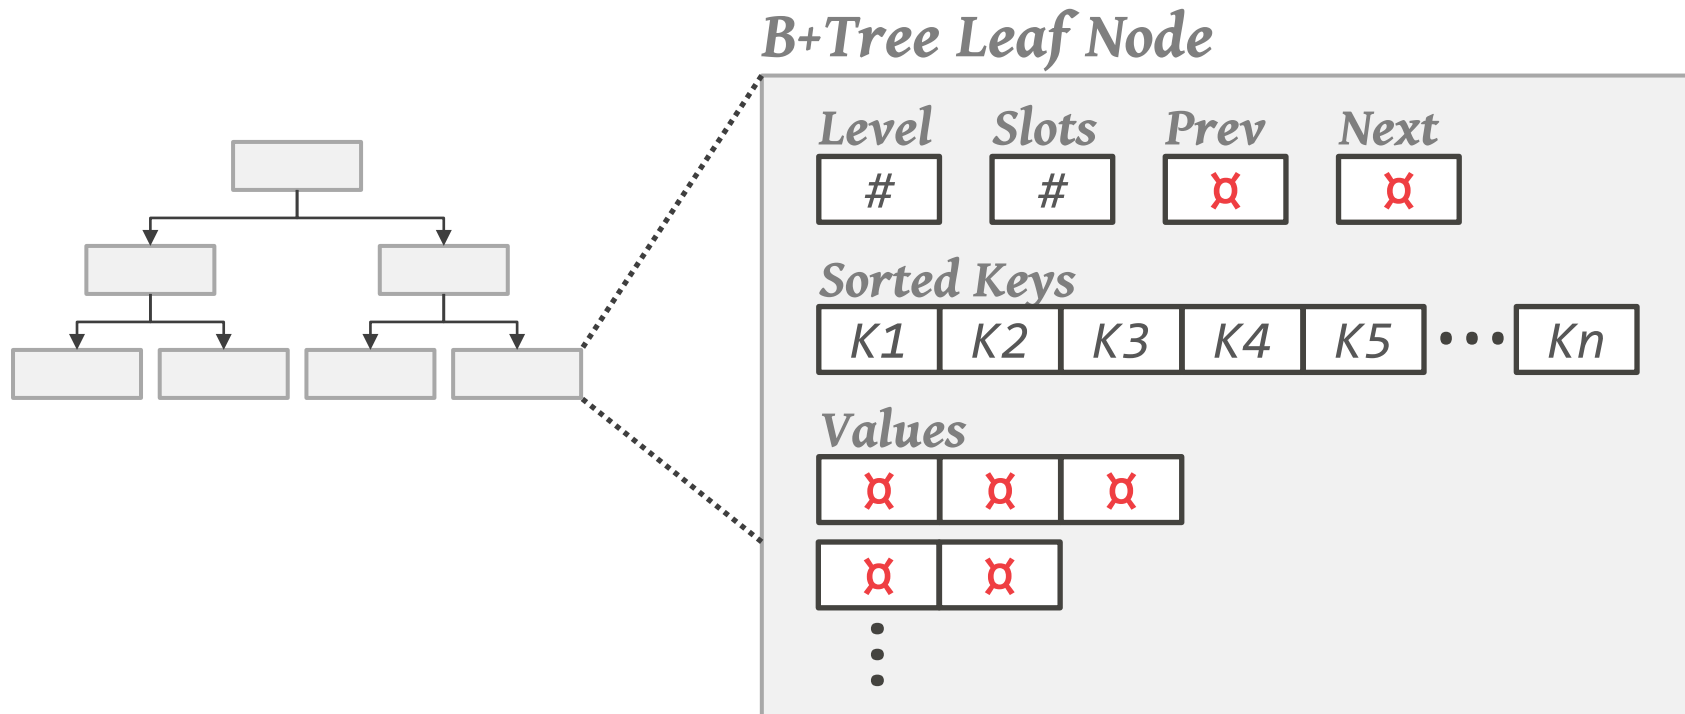
B+TREE: DUPLICATE KEYS



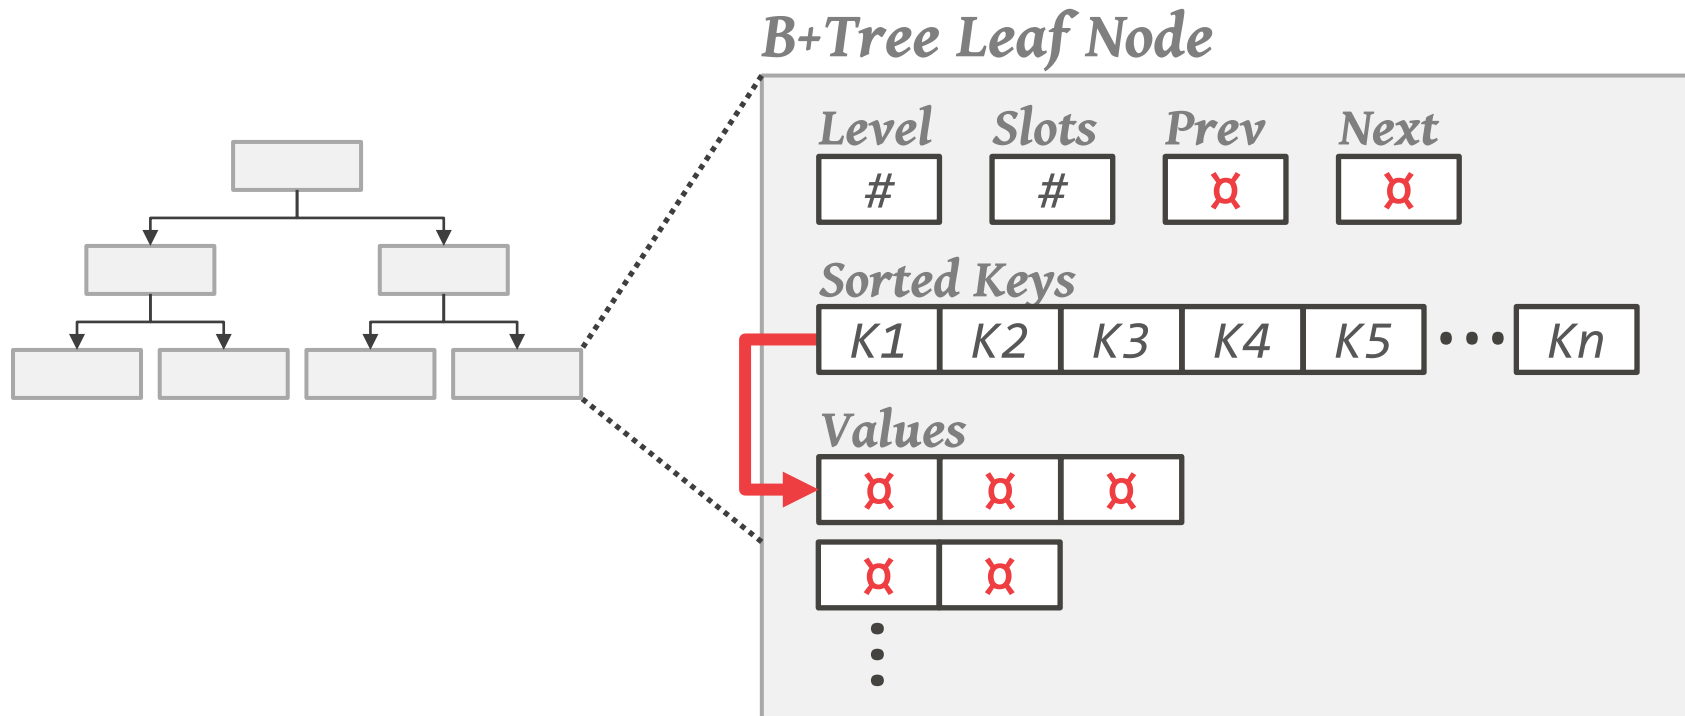
B+TREE: DUPLICATE KEYS



B+TREE: VALUE LISTS



B+TREE: VALUE LISTS



B+TREE: VARIABLE LENGTH KEYS

Approach #1: Pointers

→ Store the keys as pointers to the tuple's attribute.

Approach #2: Variable Length Nodes


→ The size of each node in the b+tree can vary.

→ Requires careful memory management.

Approach #3: Key Map

→ Embed an array of pointers that map to the key + value list within the node.

B+TREE: KEY MAP



B+Tree Leaf Node

<i>Level</i>	<i>Slots</i>	<i>Prev</i>	<i>Next</i>
#	#	⌘	⌘

Key Map

⌘	⌘	⌘
---	---	---

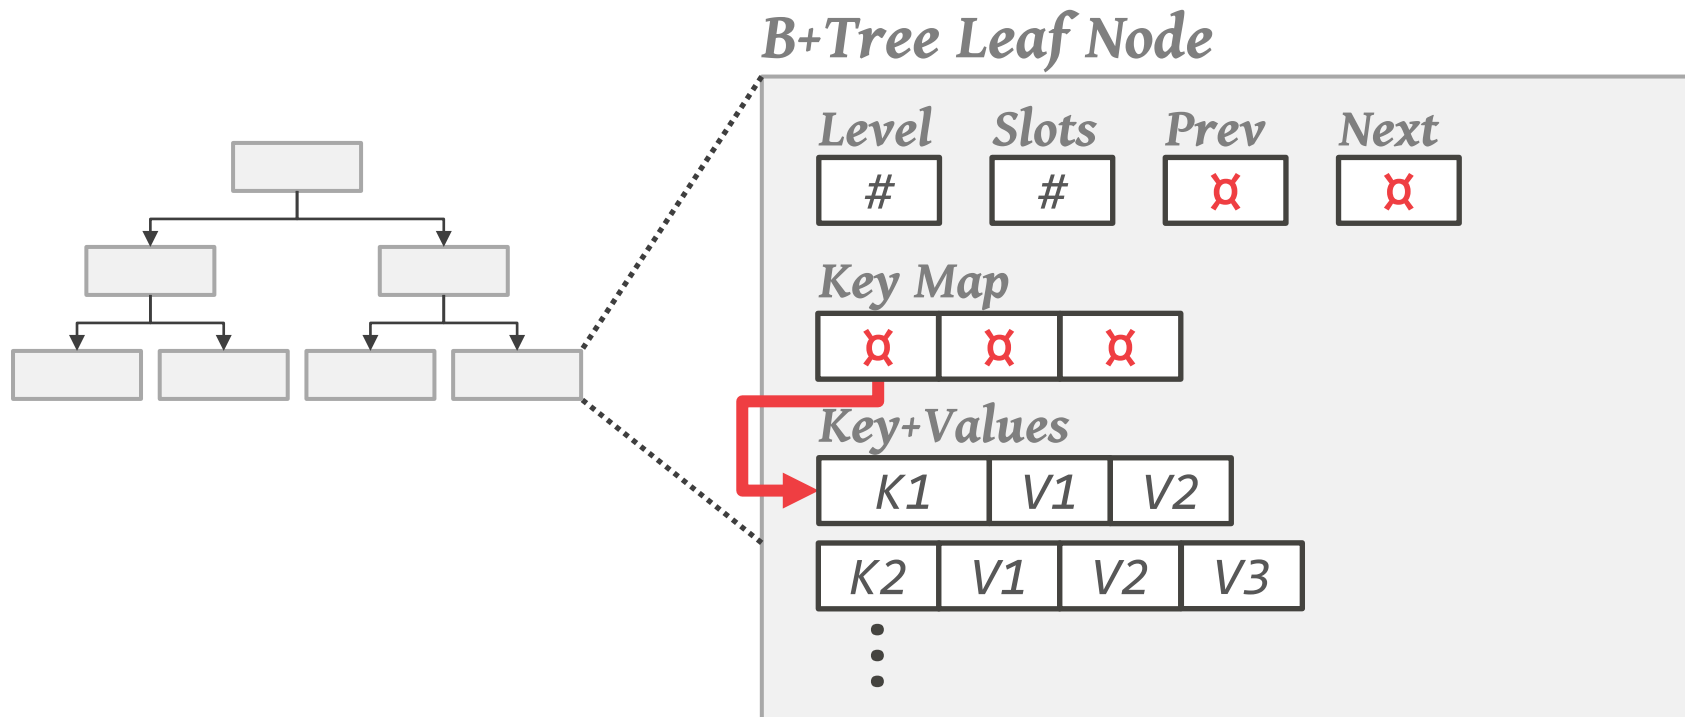
Key+Values

K1	V1	V2
----	----	----

K2	V1	V2	V3
----	----	----	----

⋮

B+TREE: KEY MAP



B+TREE ALTERNATIVES

T-Trees

Skip Lists

Radix Trees (aka Patricia Trees)

MassTree

Fractal Trees

B+TREE ALTERNATIVES

T-Trees

Skip Lists

Radix Trees (aka Patricia Trees)

MassTree

Fractal Trees

T-TREES


Based on AVL Trees. Instead of storing keys in nodes, store pointers to their original values.

Proposed in 1986 from Univ. of Wisconsin
Used in TimesTen and other early in-memory DBMSs during the 1990s.




A STUDY OF INDEX STRUCTURES FOR MAIN
MEMORY DATABASE MANAGEMENT SYSTEMS
VLDB 1986


T-TREES



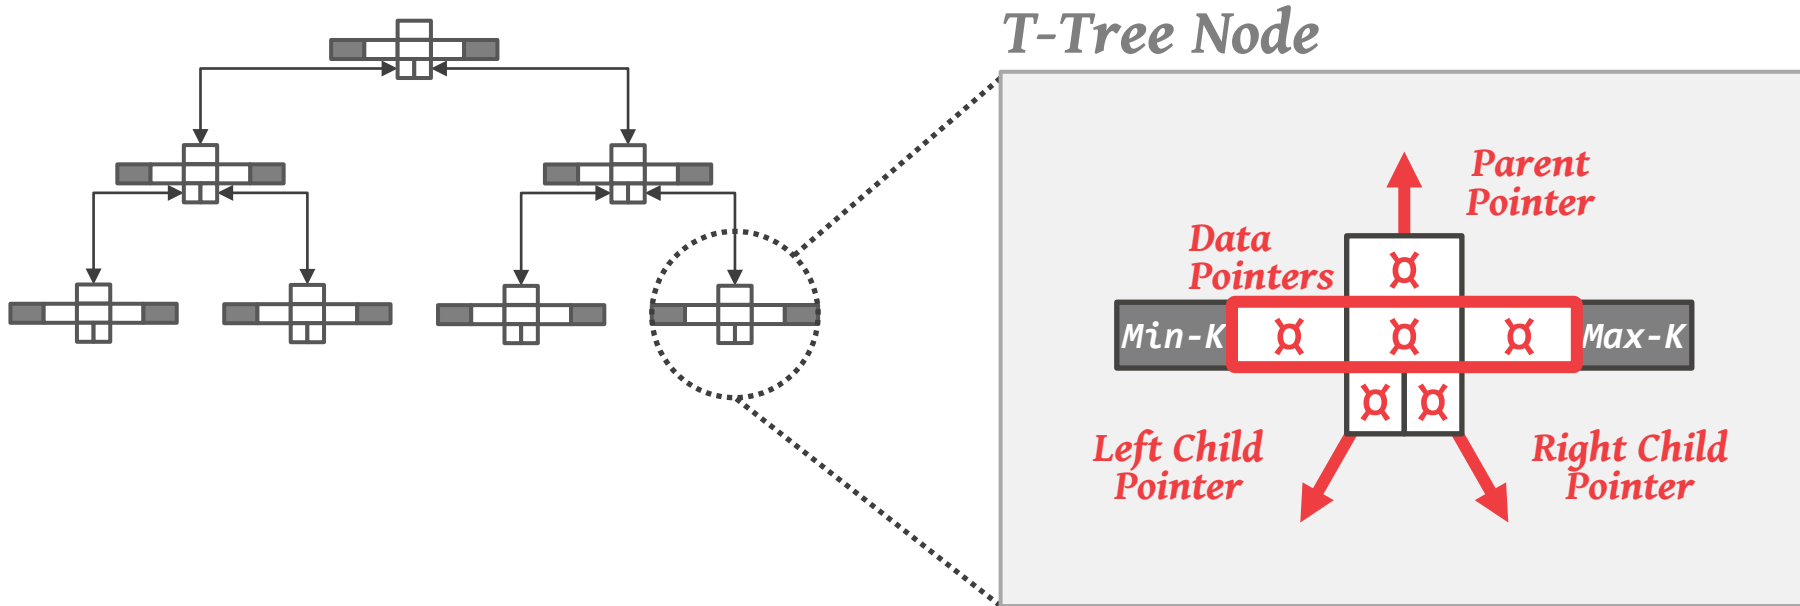
T-TREES



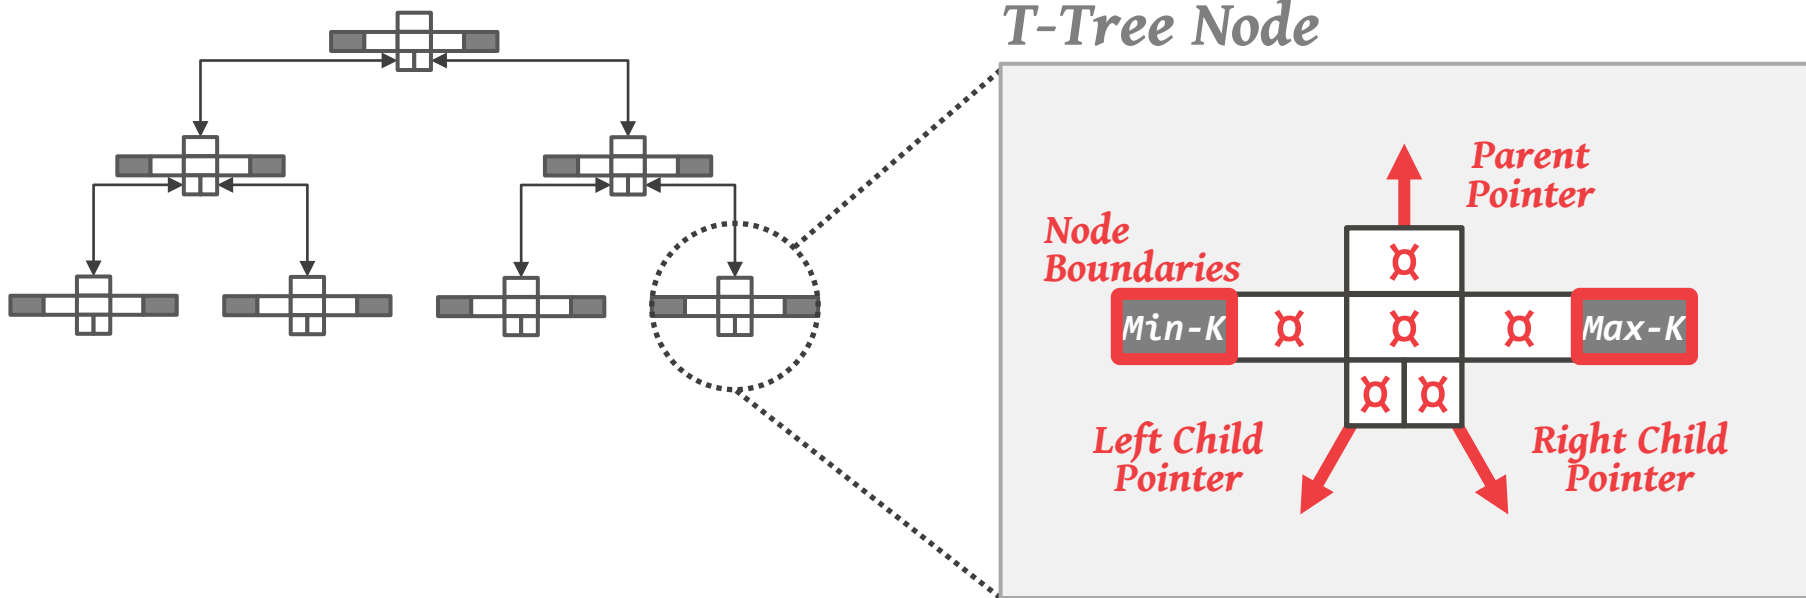
T-Tree Node



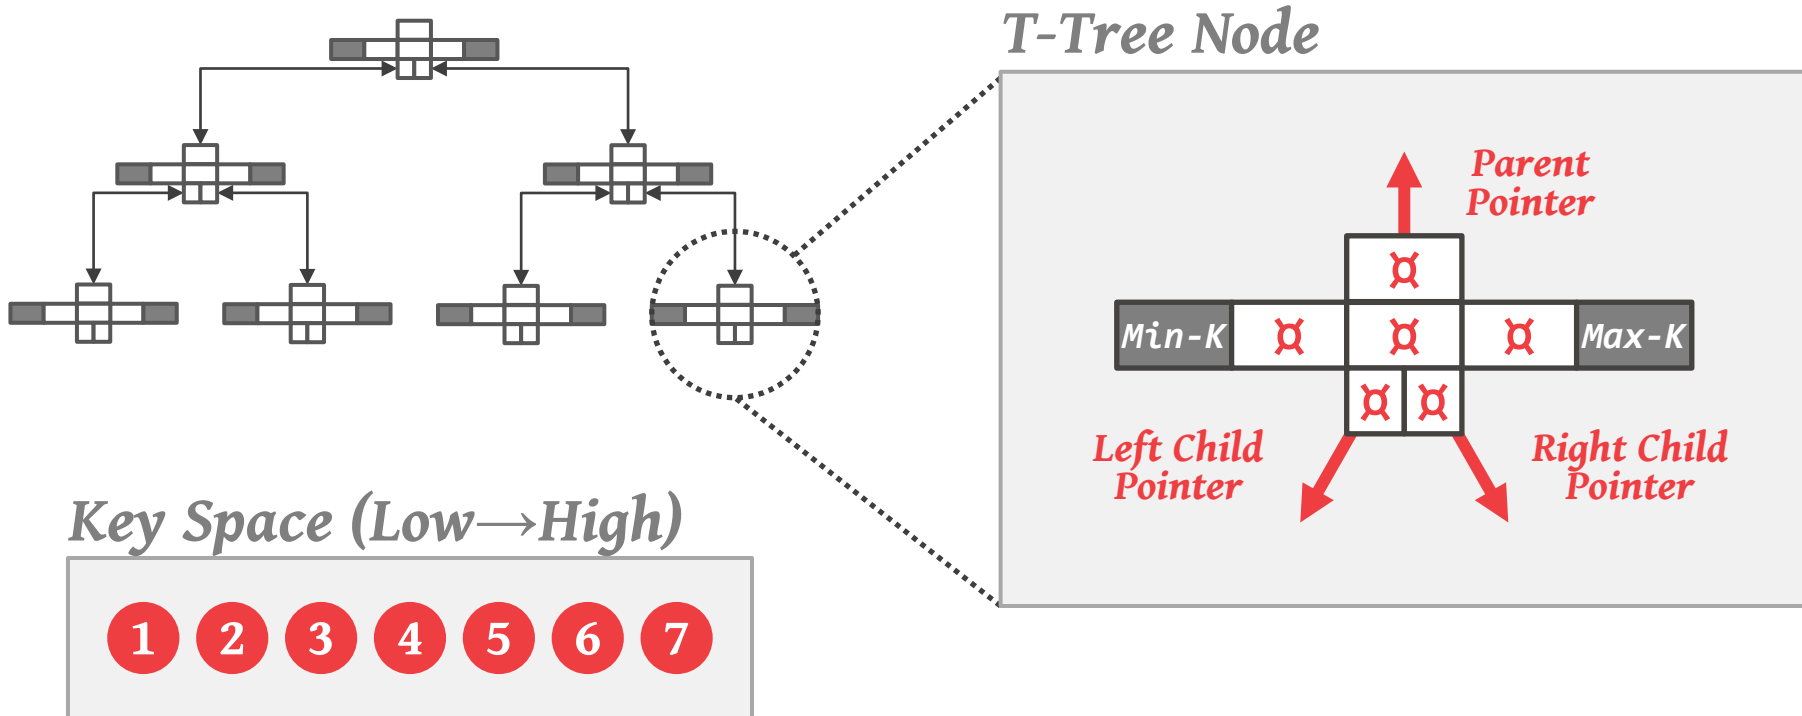
T-TREES



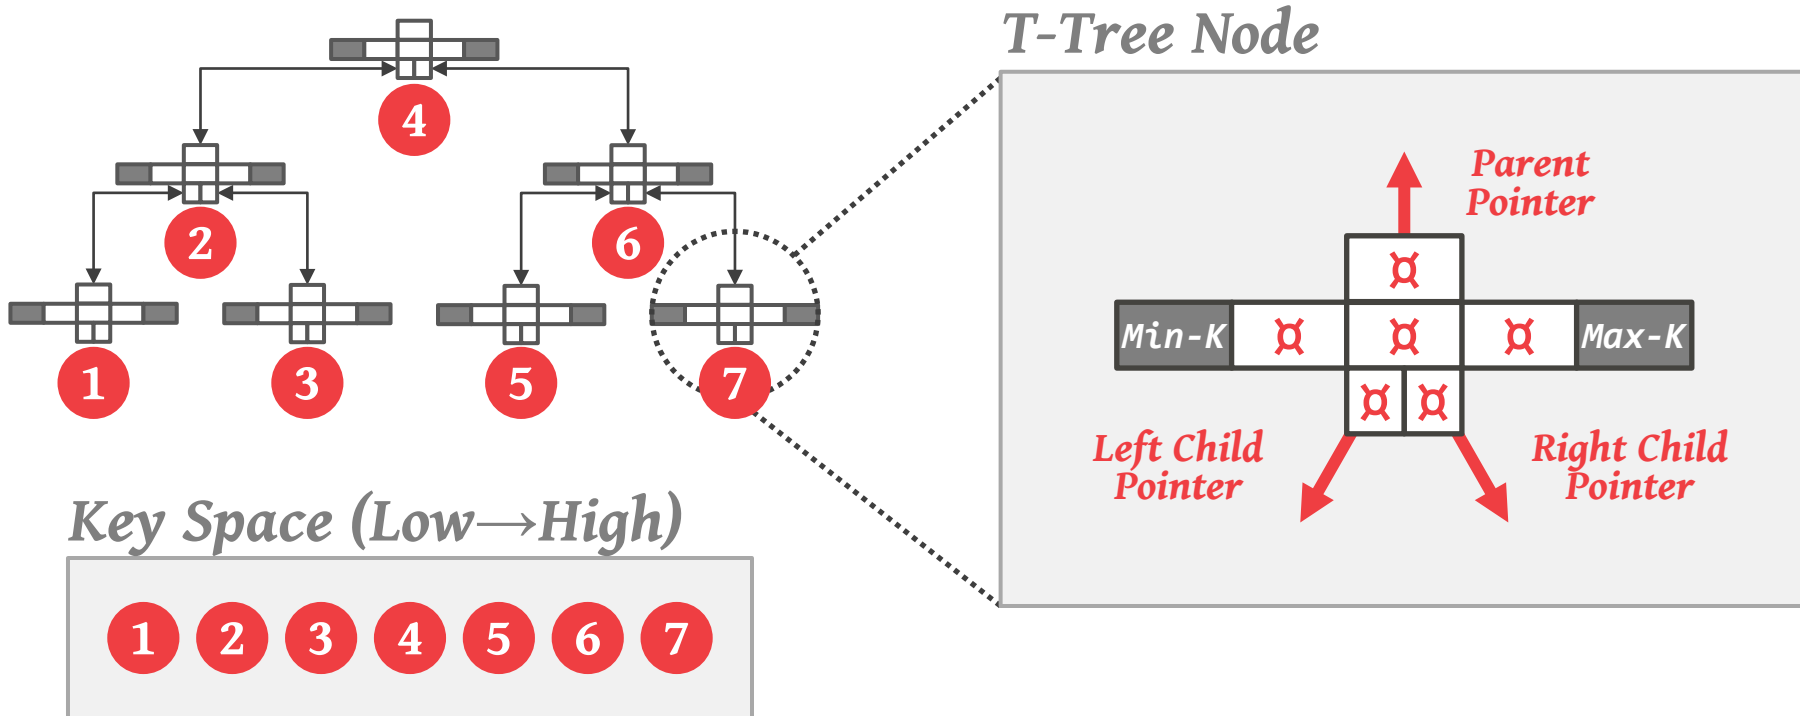
T-TREES



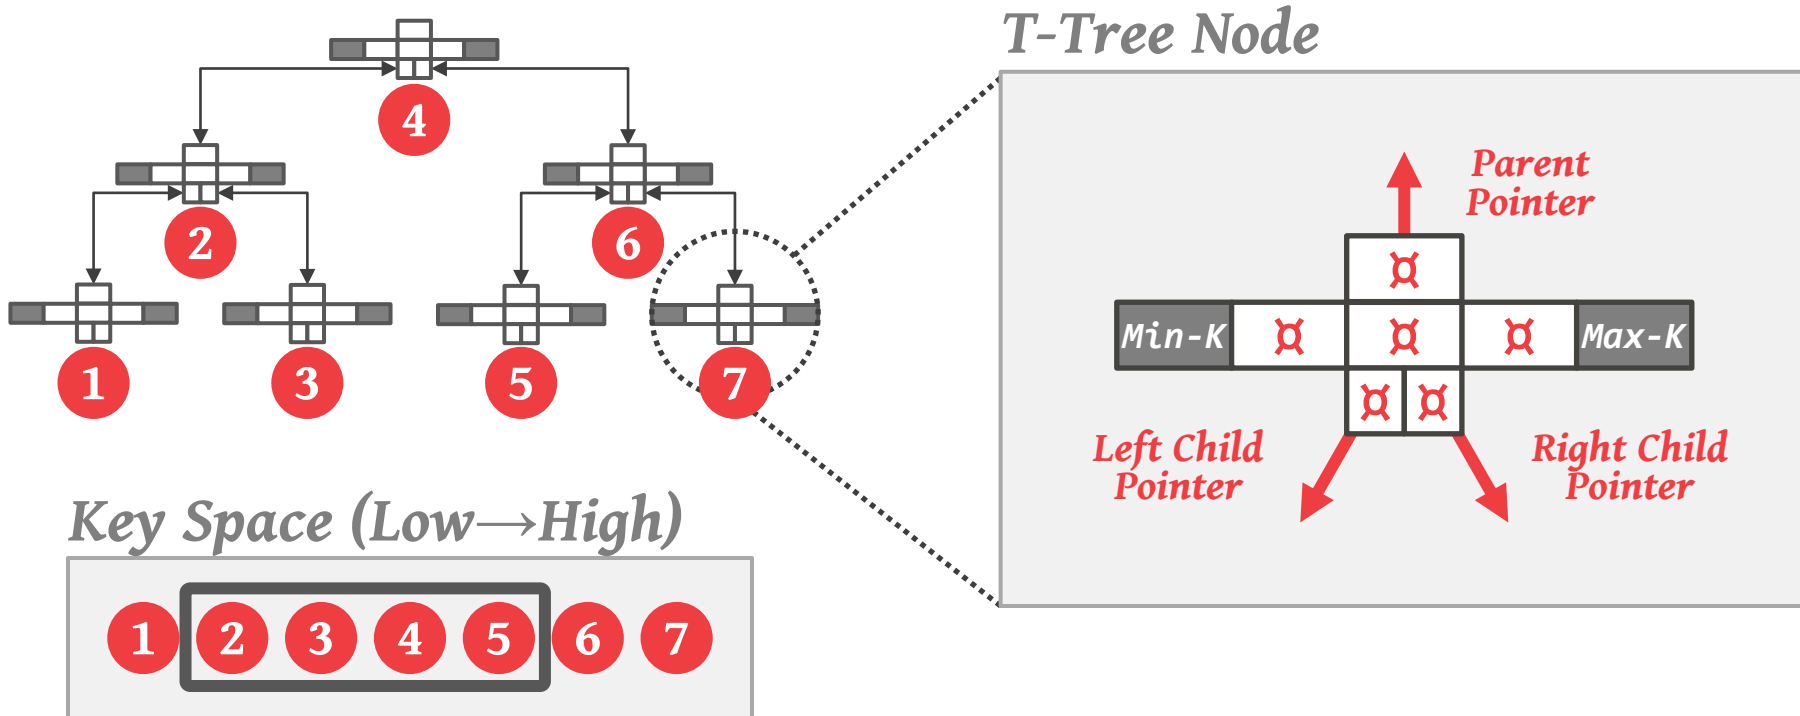
T-TREES



T-TREES



T-TREES



T-TREES

Advantages

- Uses less memory because it does not store keys inside of each node.

Disadvantages

- Have to chase pointers when scanning range or performing binary search inside of a node.
- Difficult to rebalance.
- Difficult to implement safe concurrent access.

SKIP LISTS

A collection of lists at different levels

- Lowest level is a sorted, singly linked list of all keys
- 2nd level links every other key
- 3rd level links every fourth key
- In general, a level has half the keys of one below it

To insert a new key, flip a coin to decide how many levels to add the new key into.

Provides approximate $O(\log n)$ search times.

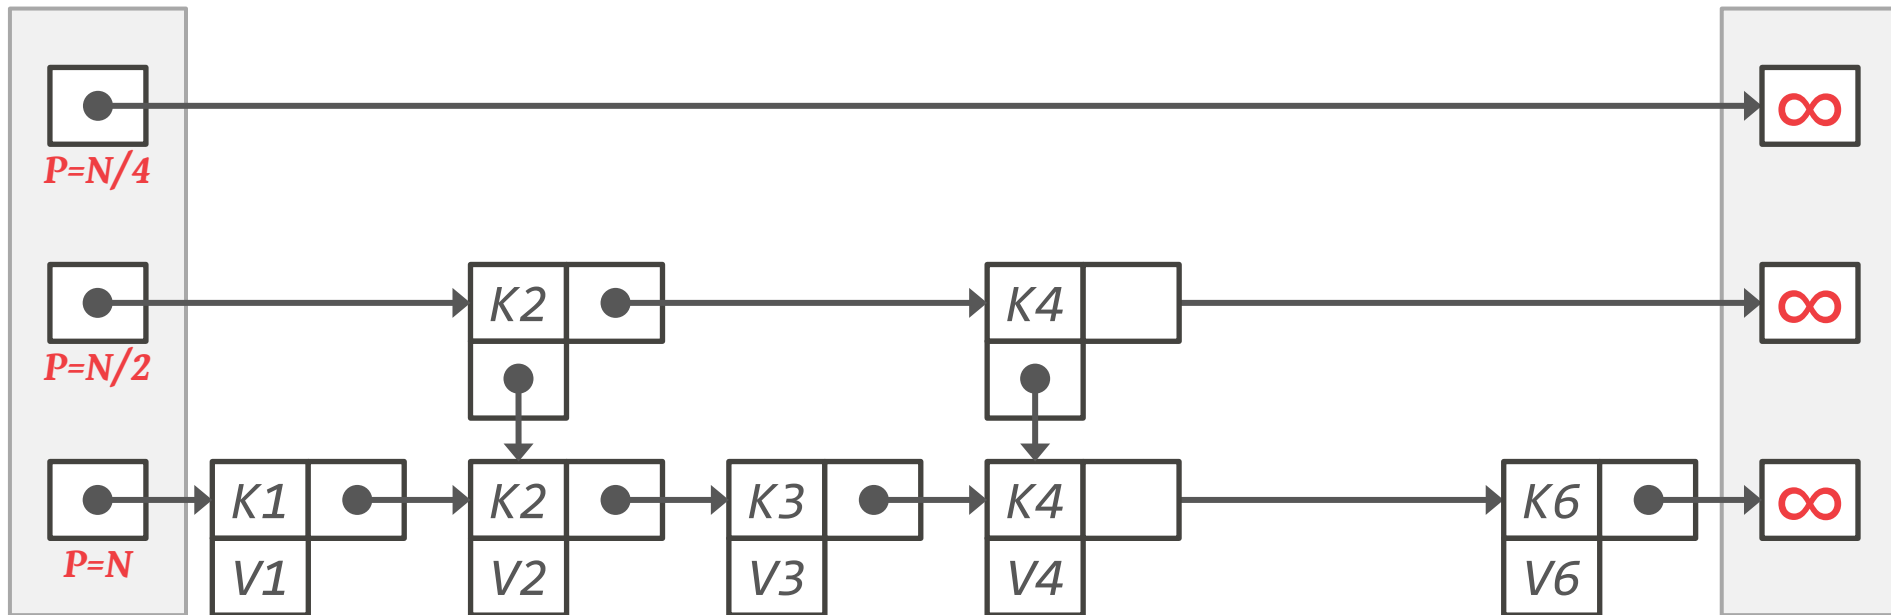


CONCURRENT MAINTENANCE OF SKIP LISTS
Univ. of Maryland Tech Report 1990

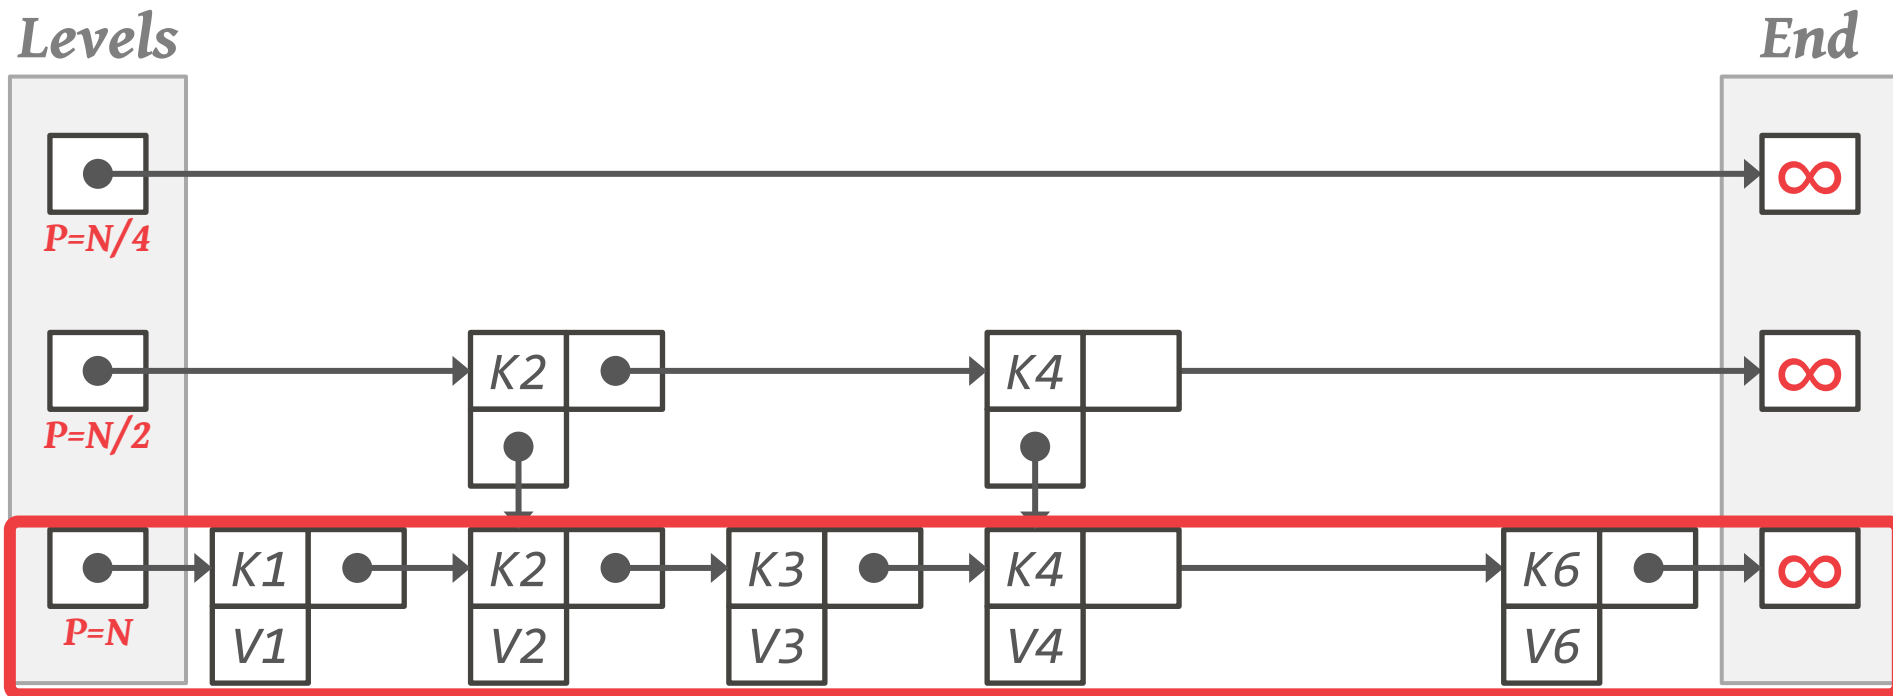
SKIP LISTS: INSERT

Levels

End



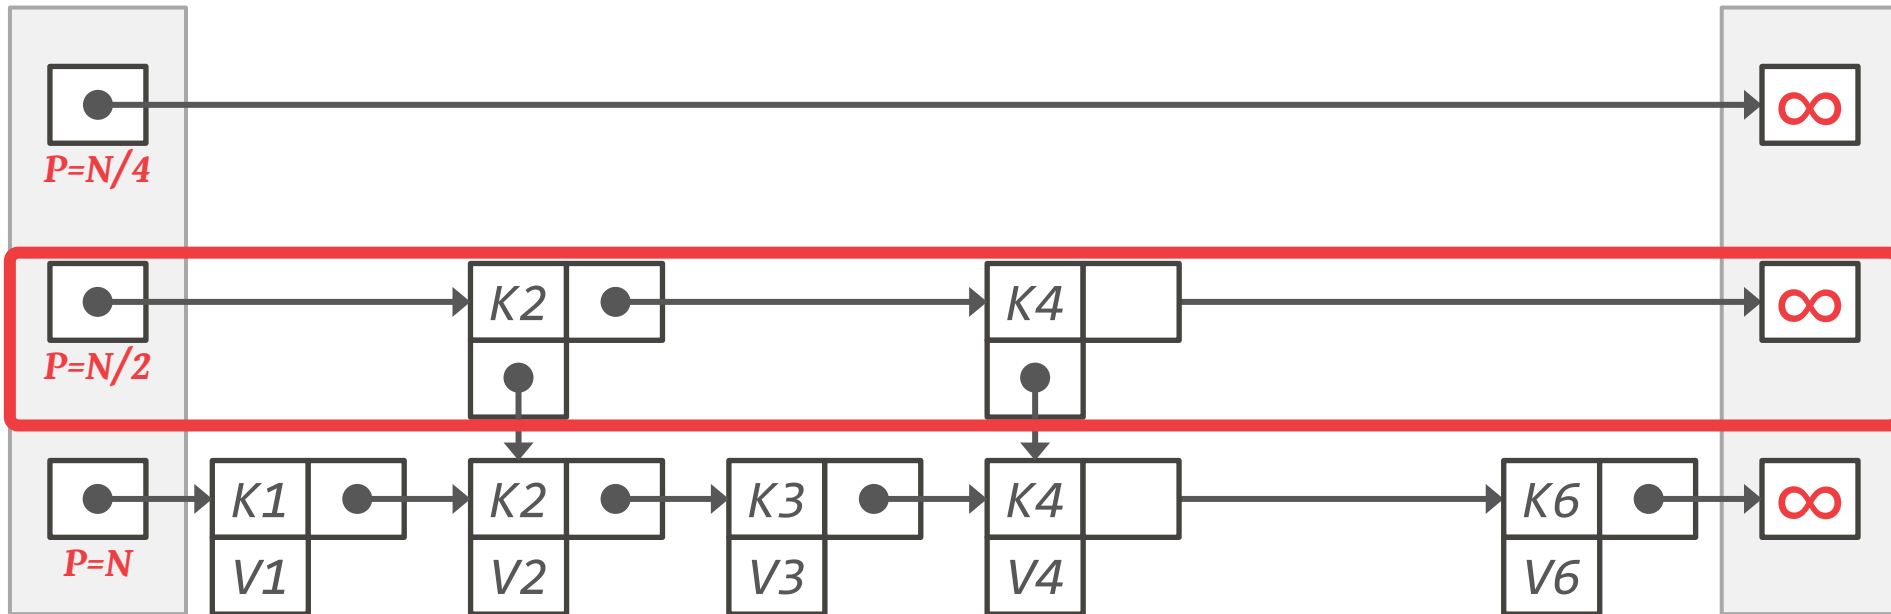
SKIP LISTS: INSERT



SKIP LISTS: INSERT

Levels

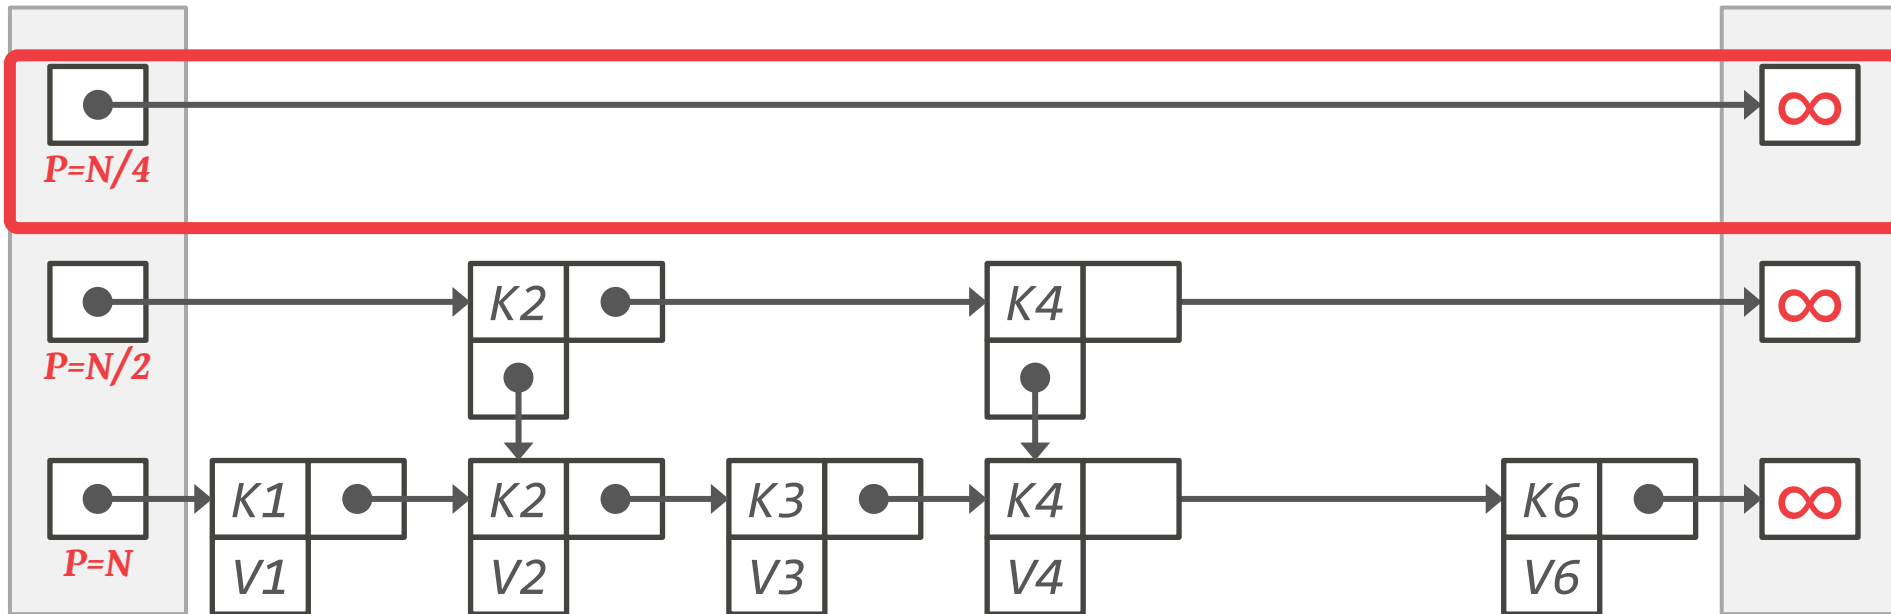
End



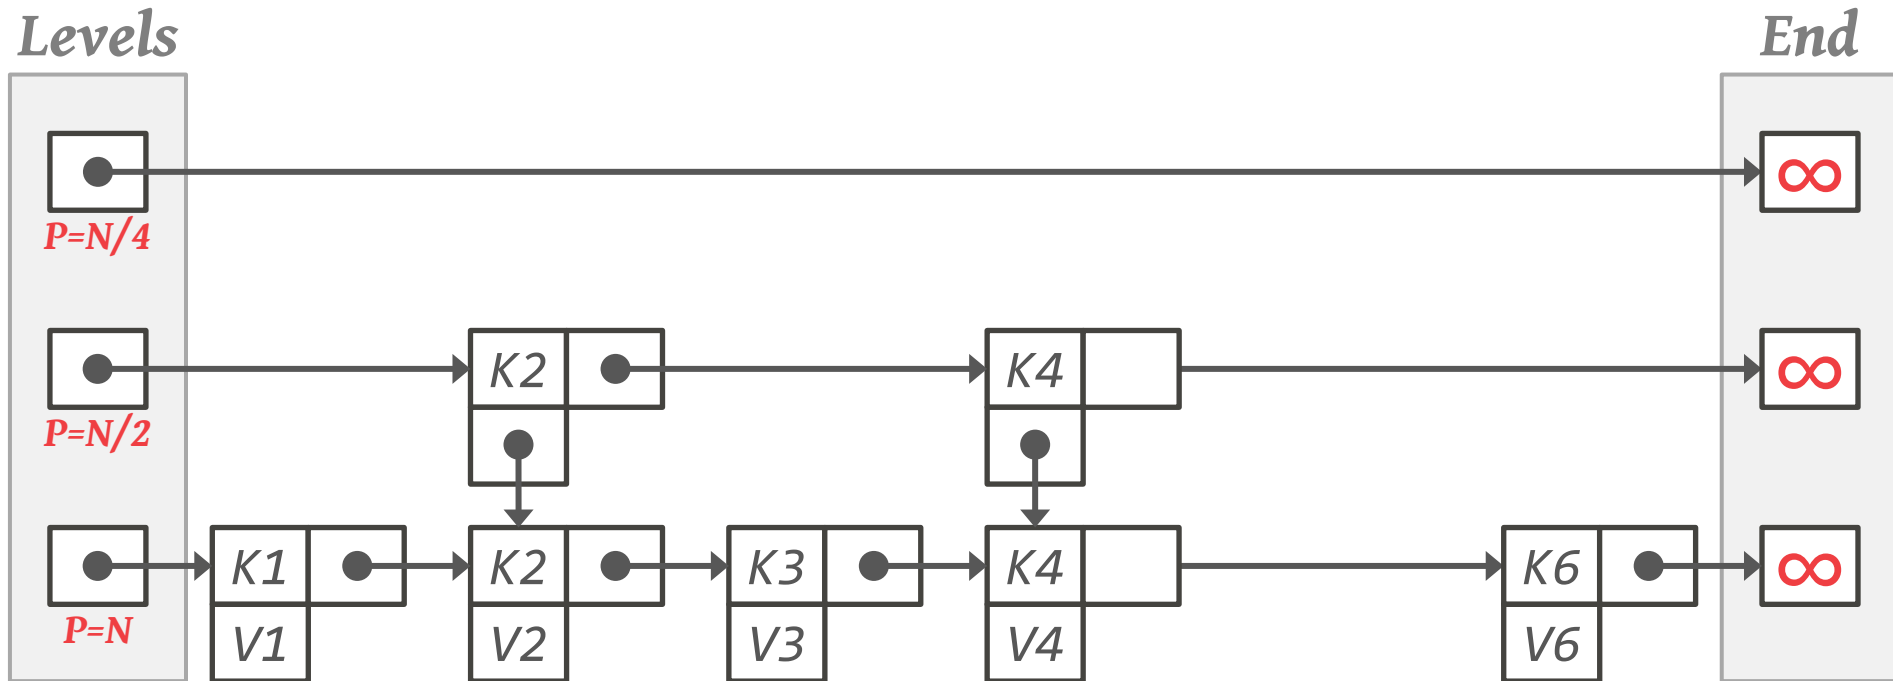
SKIP LISTS: INSERT

Levels

End



SKIP LISTS: INSERT

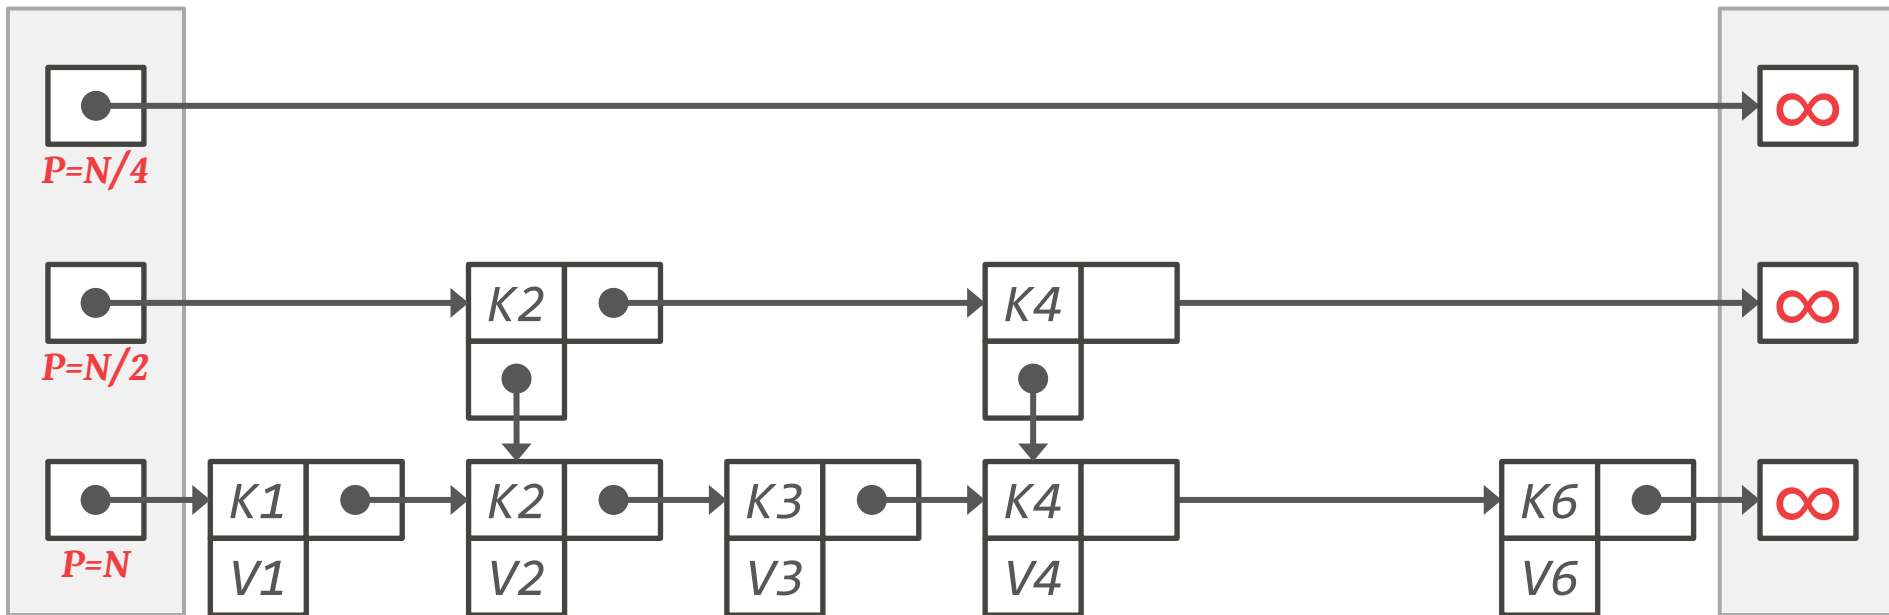


SKIP LISTS: INSERT

Txn #1: Insert K5

Levels

End

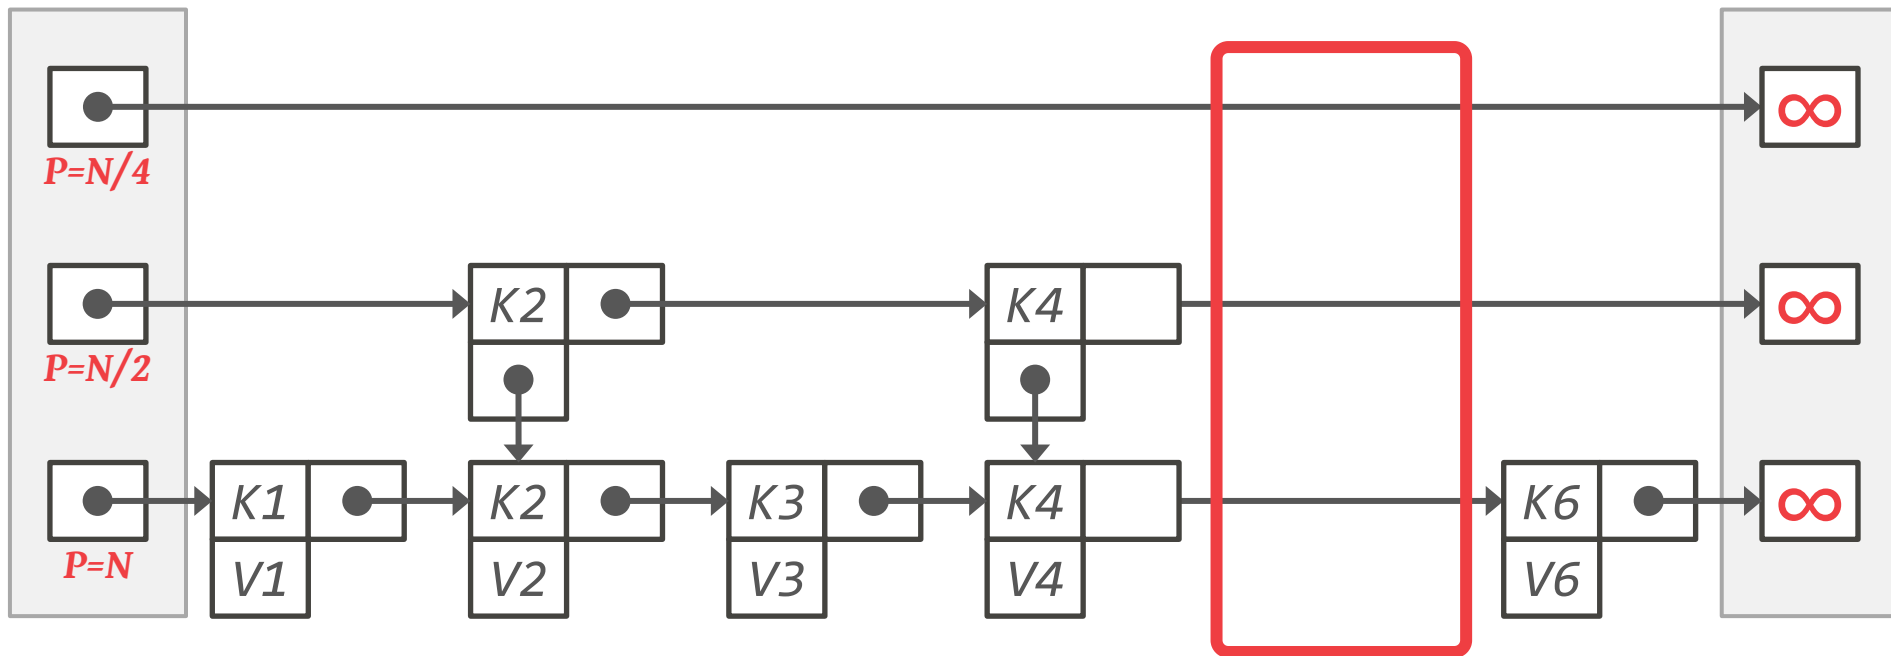


SKIP LISTS: INSERT

Txn #1: Insert K5

Levels

End

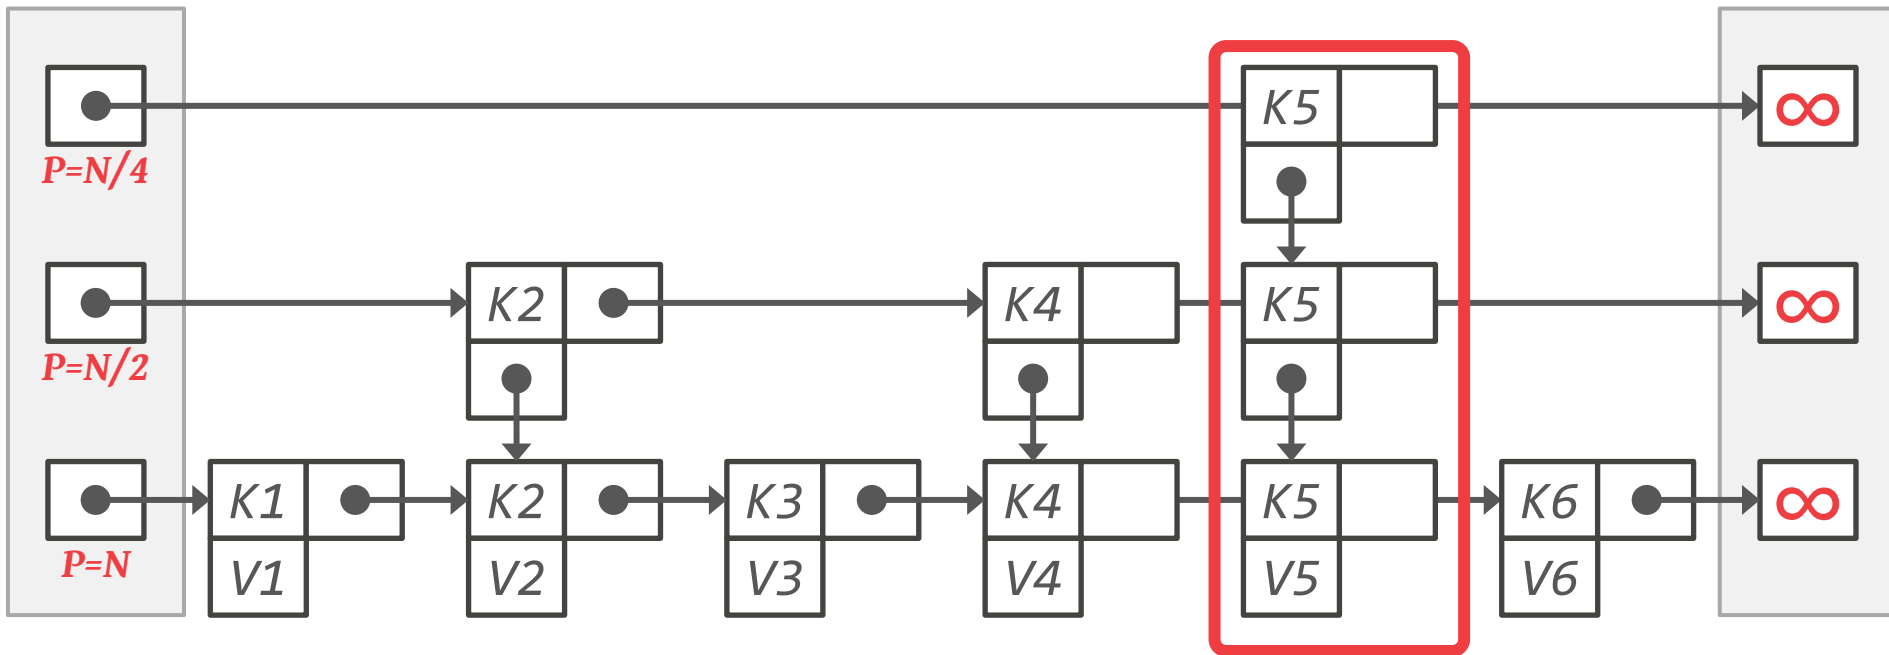


SKIP LISTS: INSERT

Txn #1: Insert K5

Levels

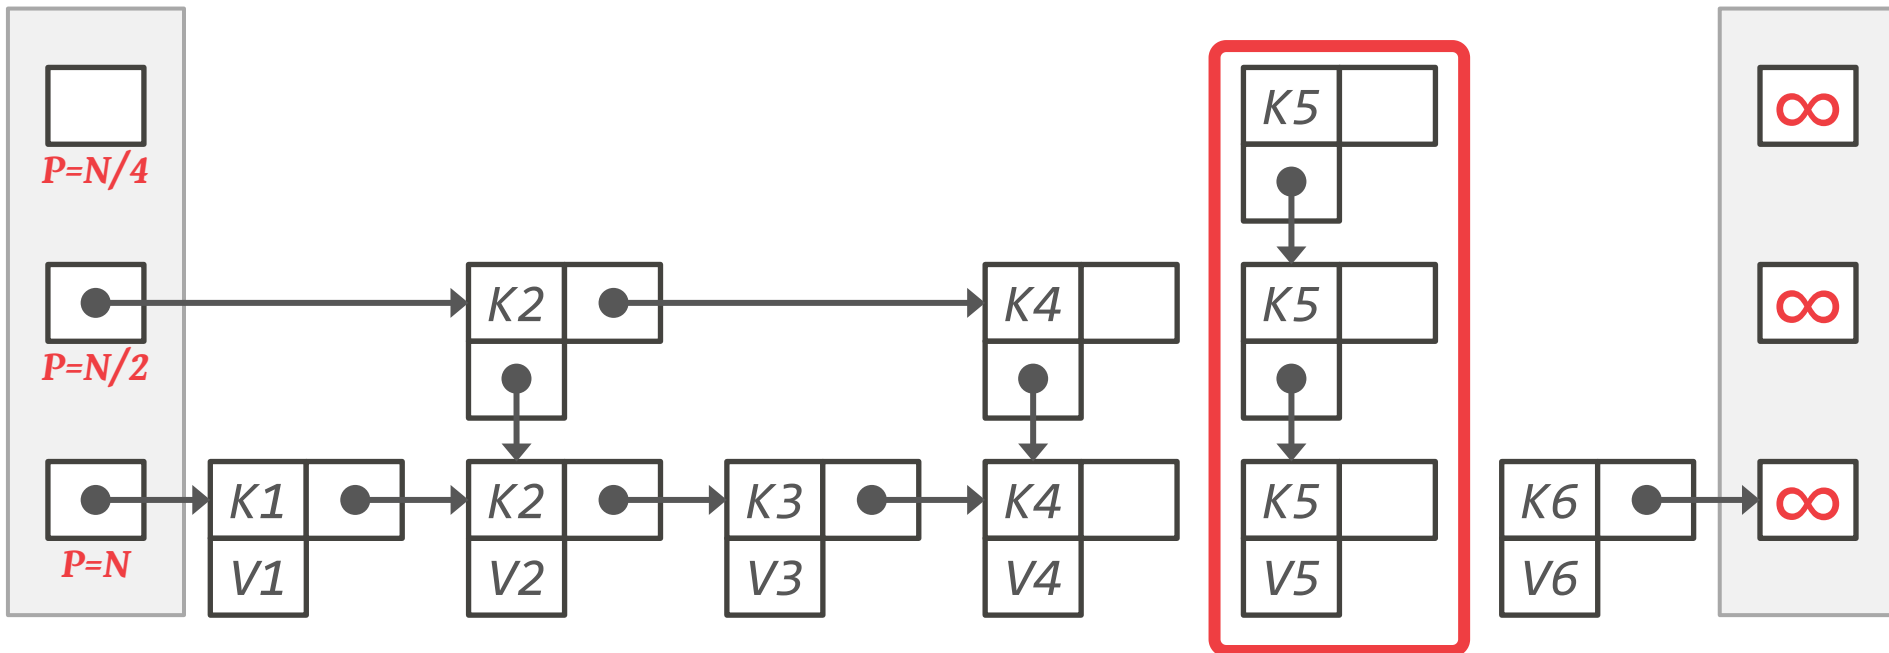
End



SKIP LISTS: INSERT

Txn #1: Insert K5

Levels

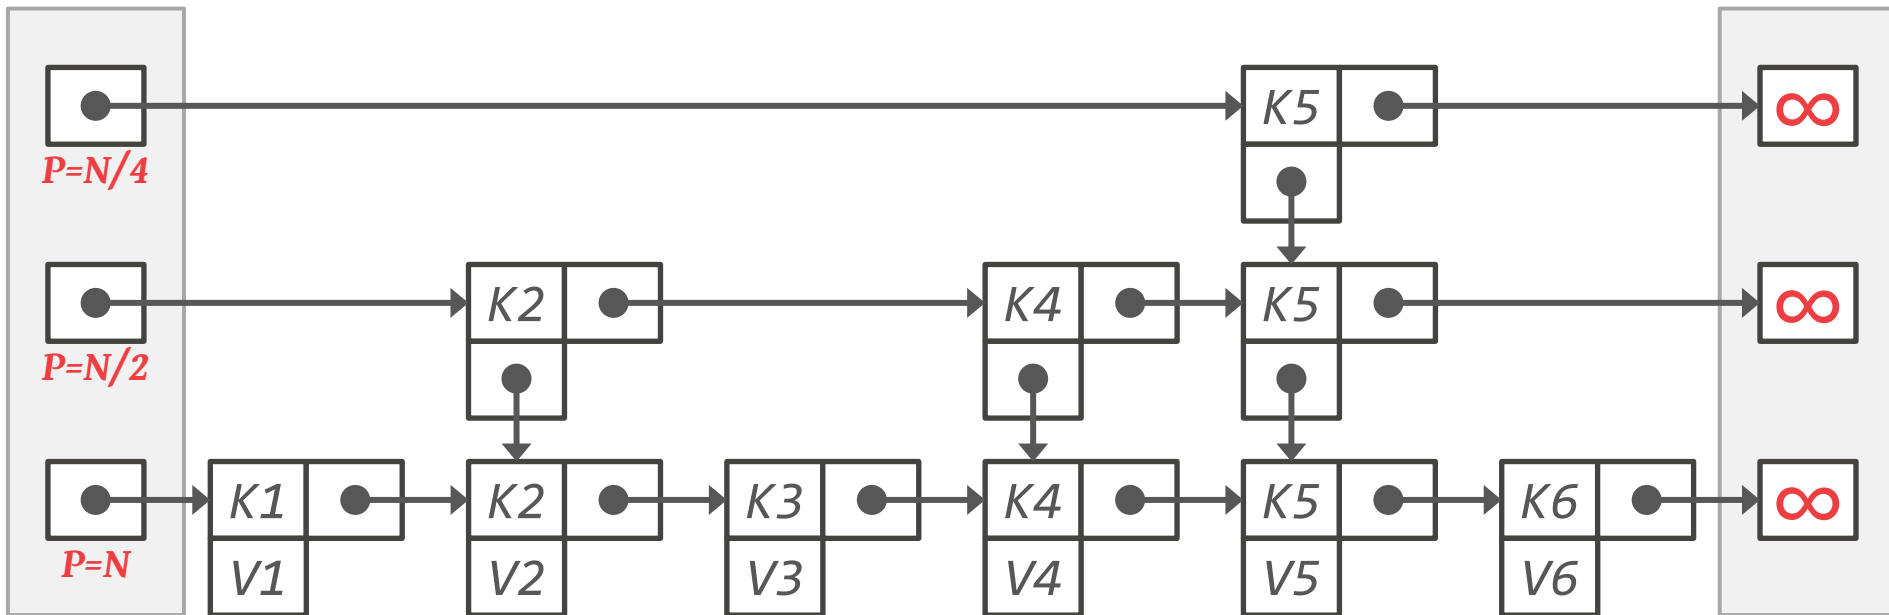


SKIP LISTS: INSERT

Txn #1: Insert K5

Levels

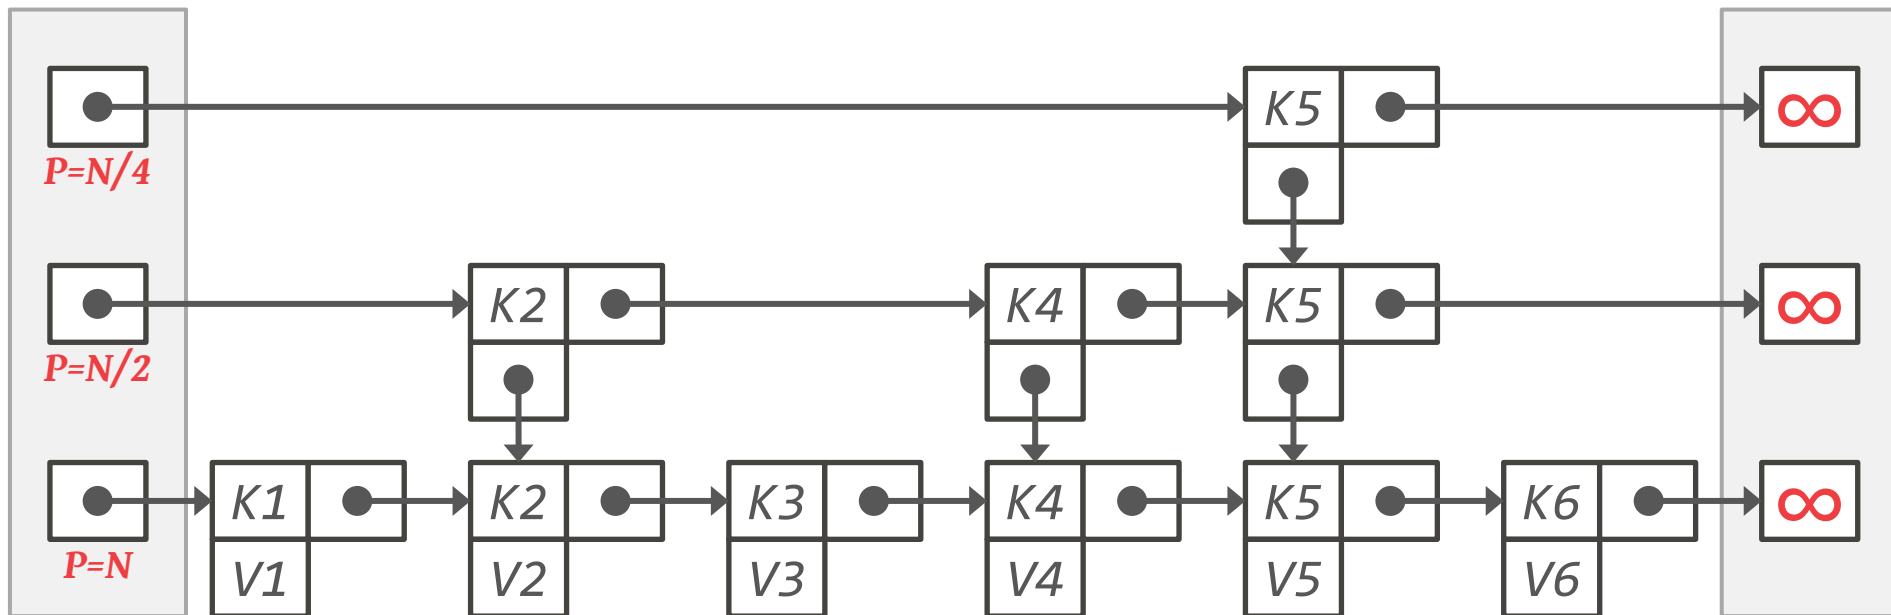
End



SKIP LISTS: SEARCH

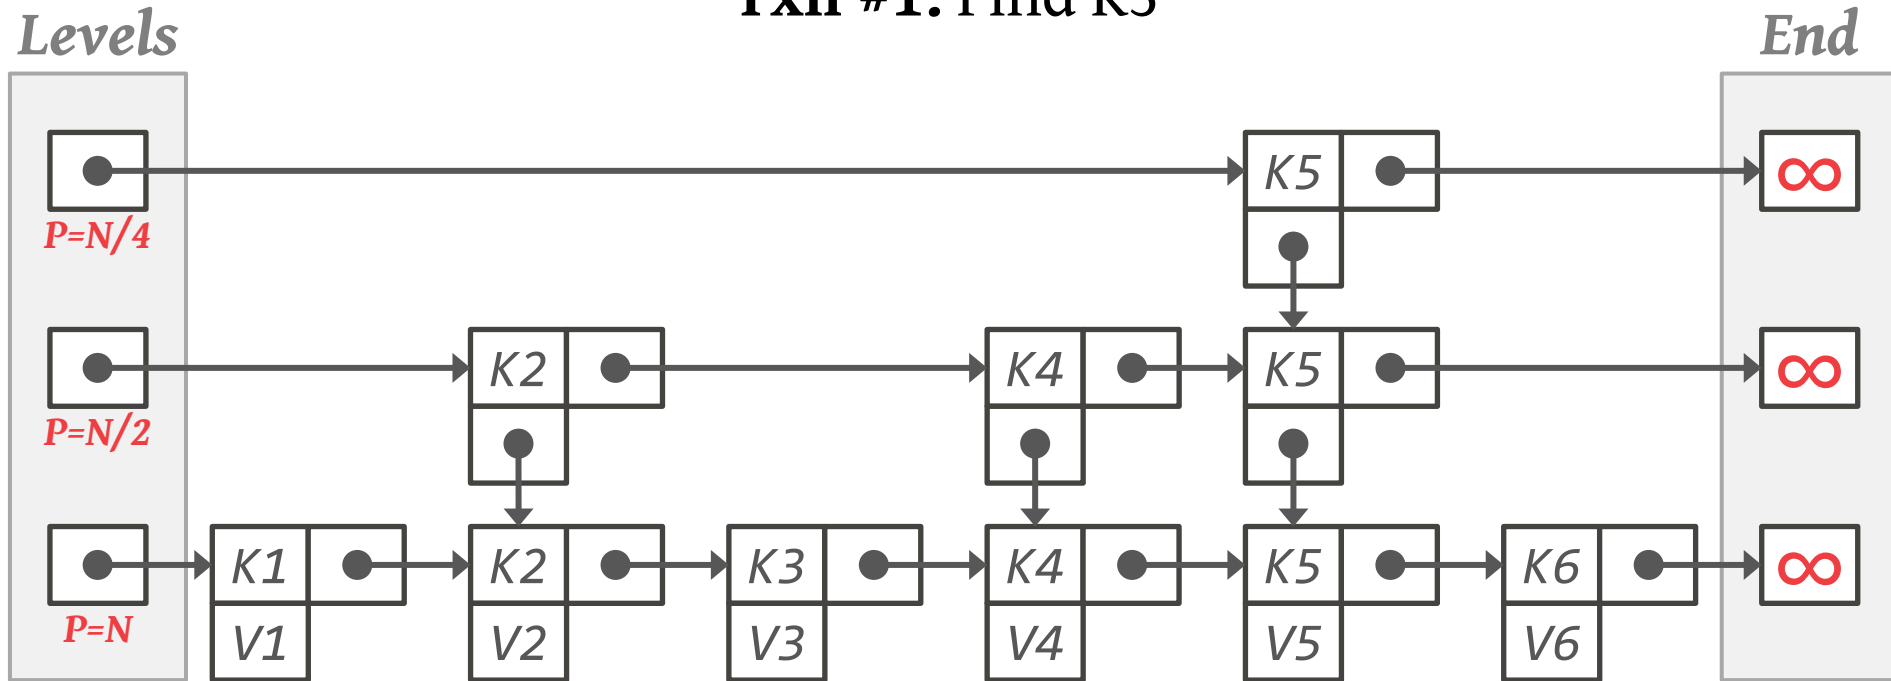
Levels

End



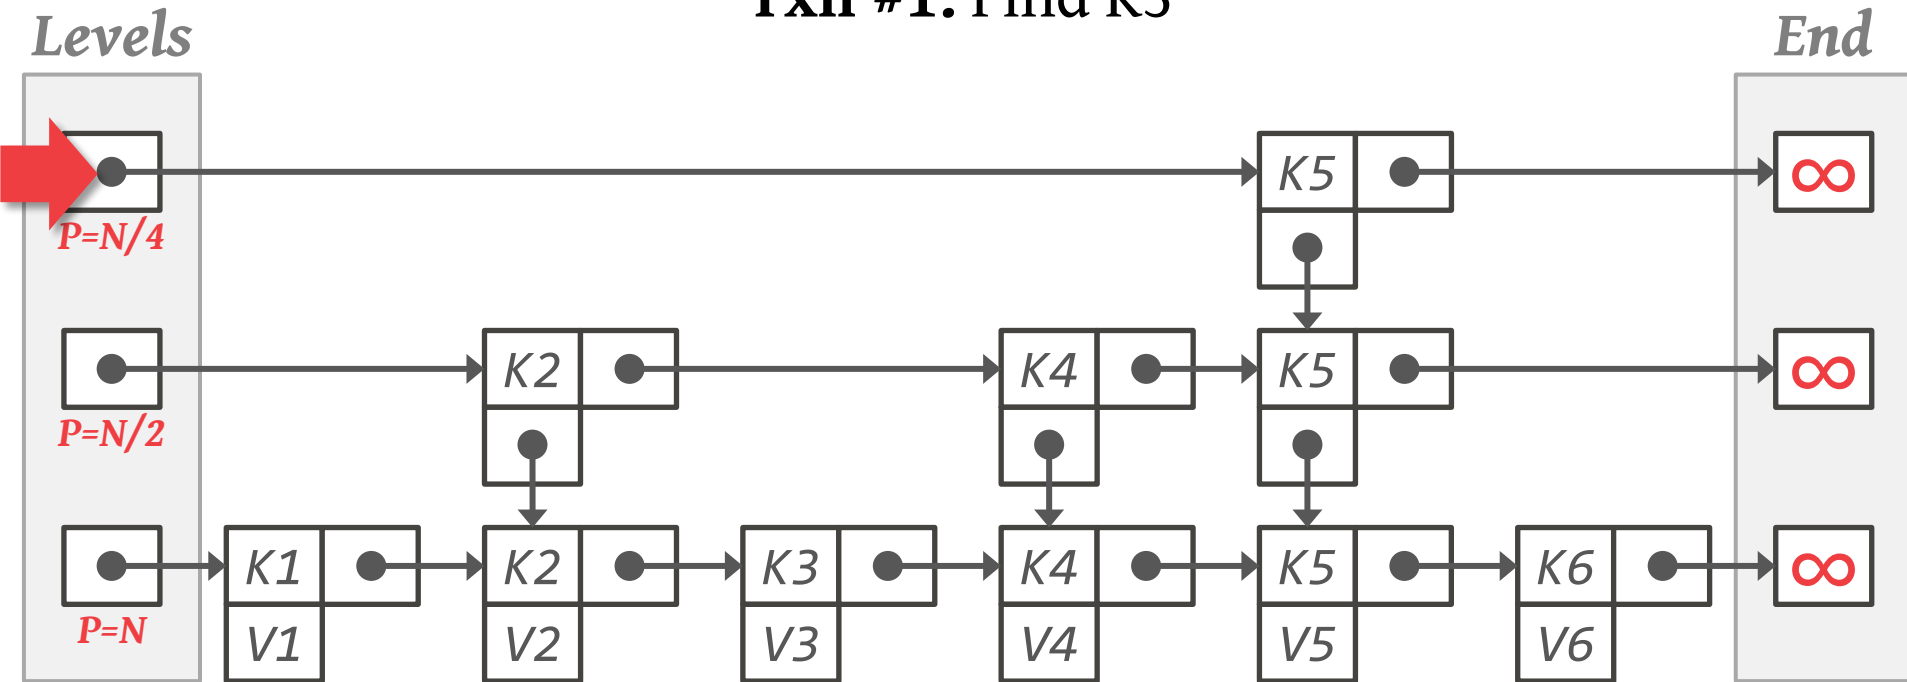
SKIP LISTS: SEARCH

Txn #1: Find K3



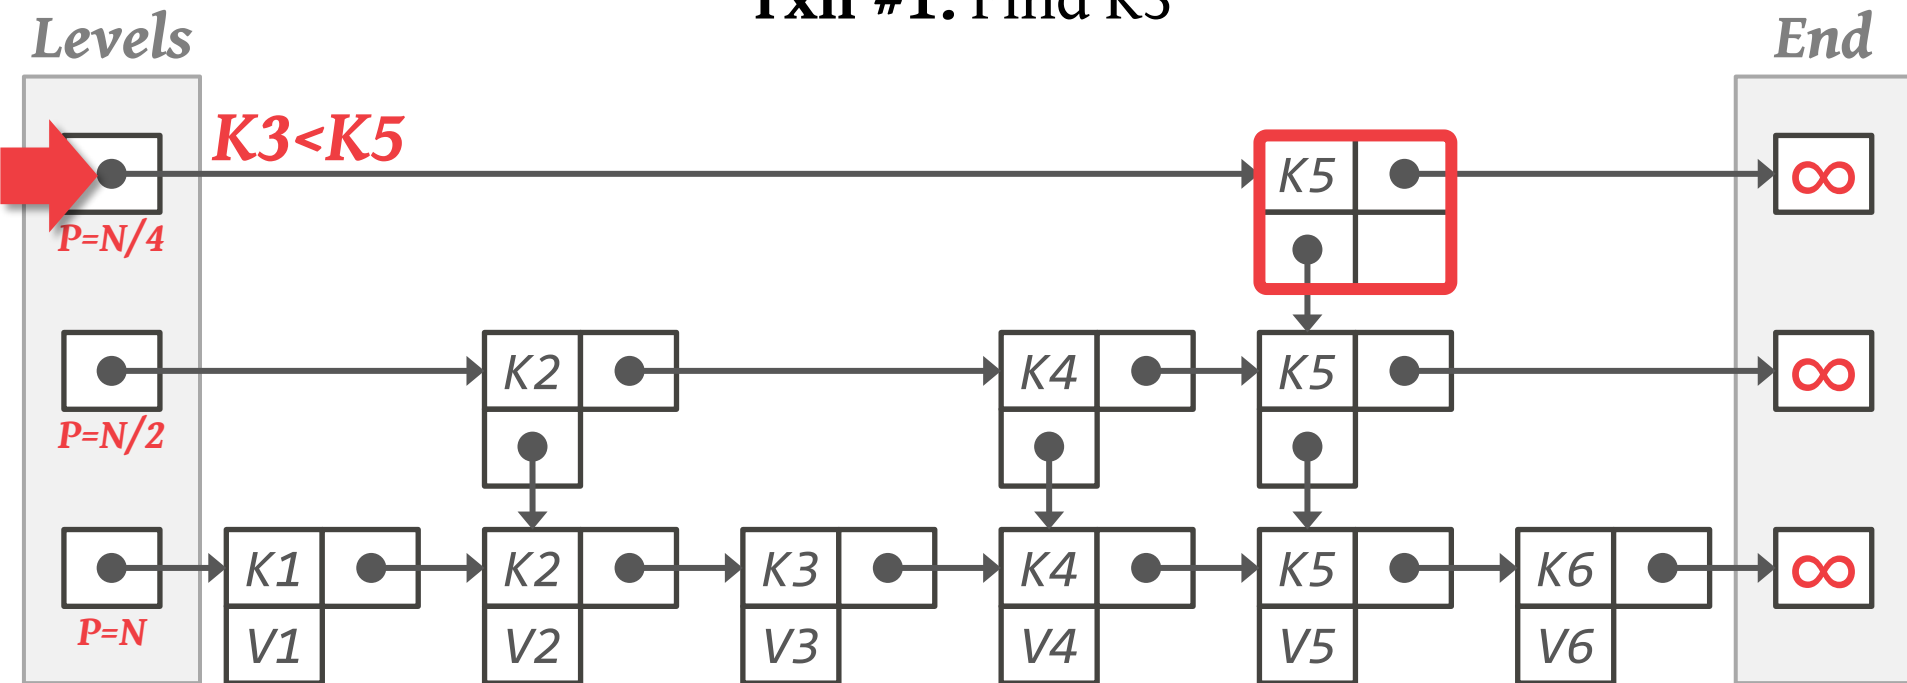
SKIP LISTS: SEARCH

Txn #1: Find K3



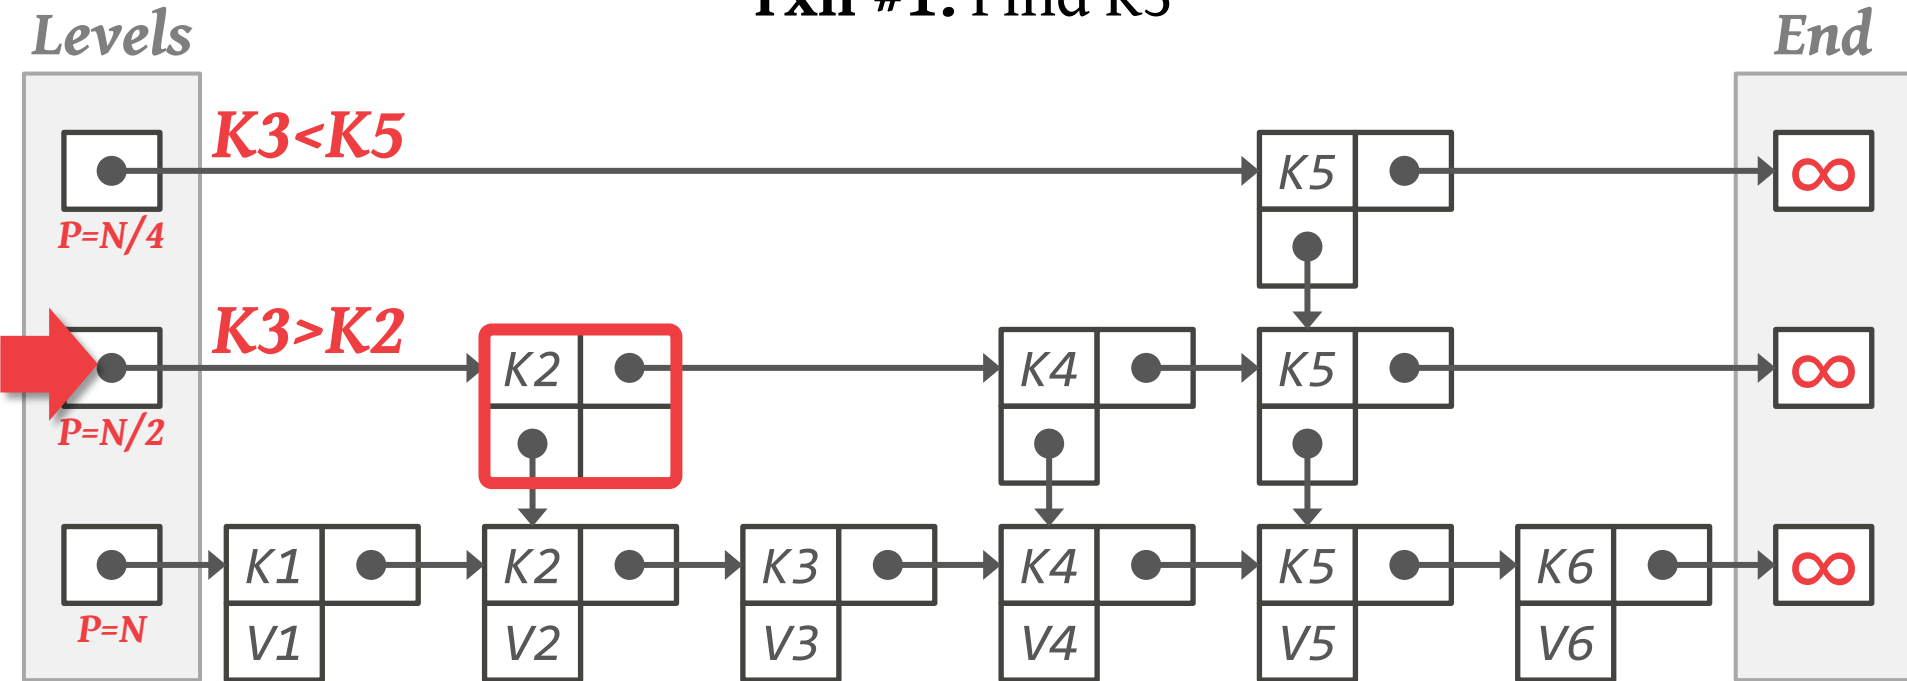
SKIP LISTS: SEARCH

Txn #1: Find K3



SKIP LISTS: SEARCH

Txn #1: Find K3

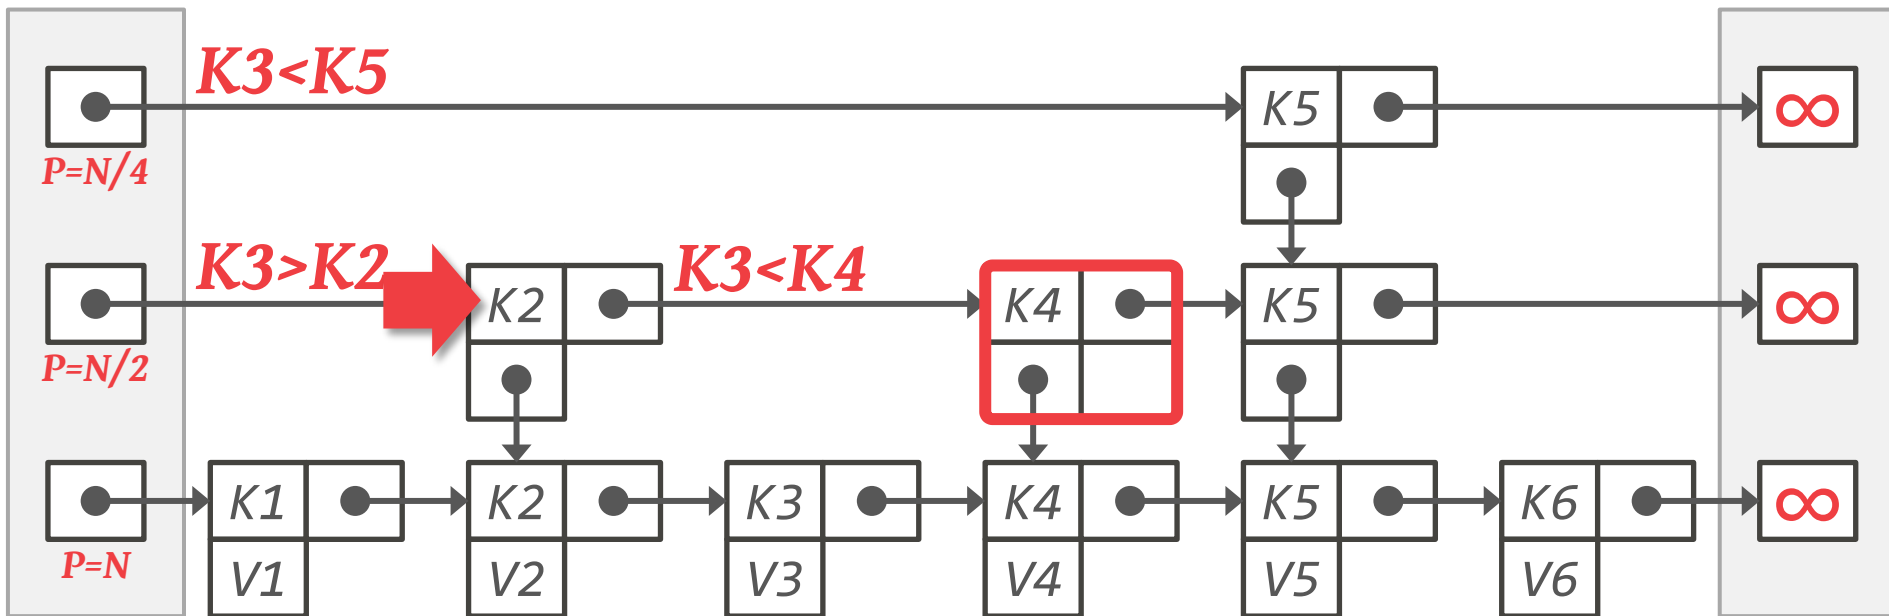


SKIP LISTS: SEARCH

Txn #1: Find K3

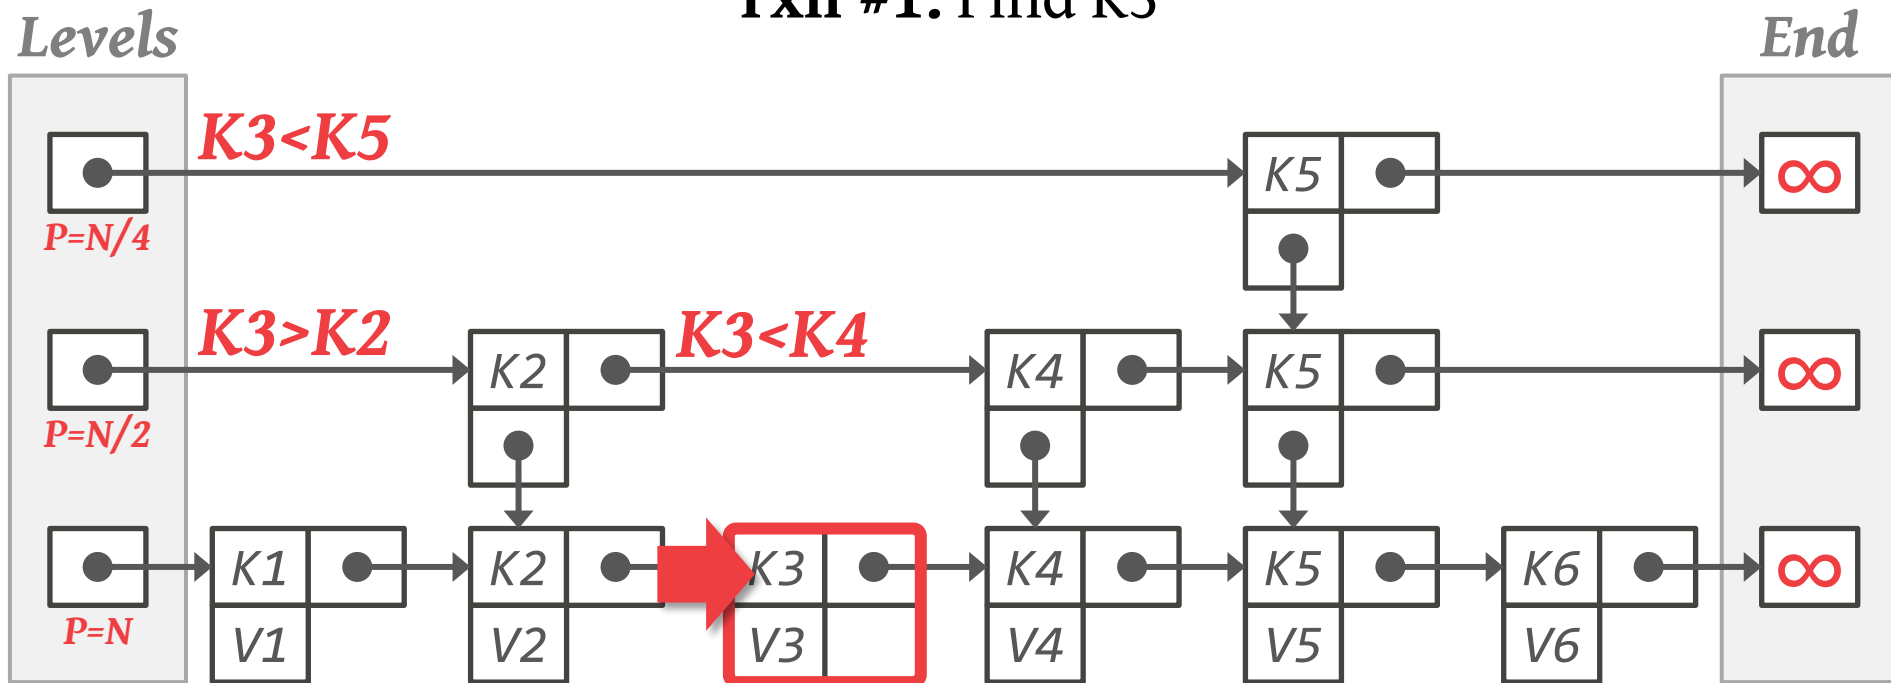
Levels

End



SKIP LISTS: SEARCH

Txn #1: Find K3



SKIP LISTS

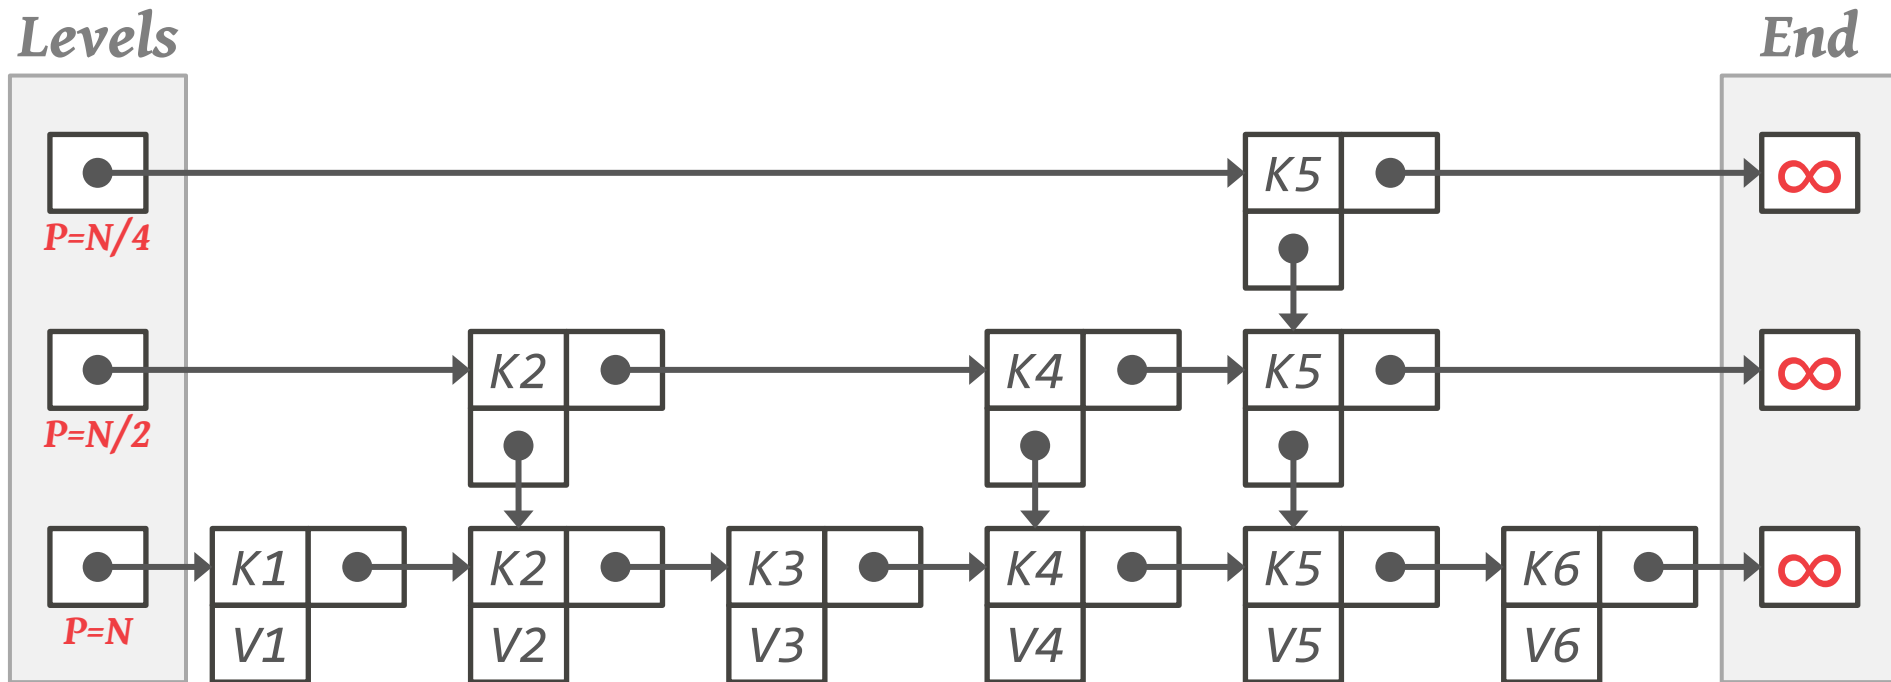
Advantages

- Uses less memory than a typical B+tree (only if you don't include reverse pointers).
- Insertions and deletions do not require rebalancing.
- It is possible to implement a concurrent skip list using only CAS instructions.

Disadvantages

- Not cache friendly because they do not optimize locality of references.
- Reverse search is non-trivial.

SKIP LISTS: REVERSE SEARCH

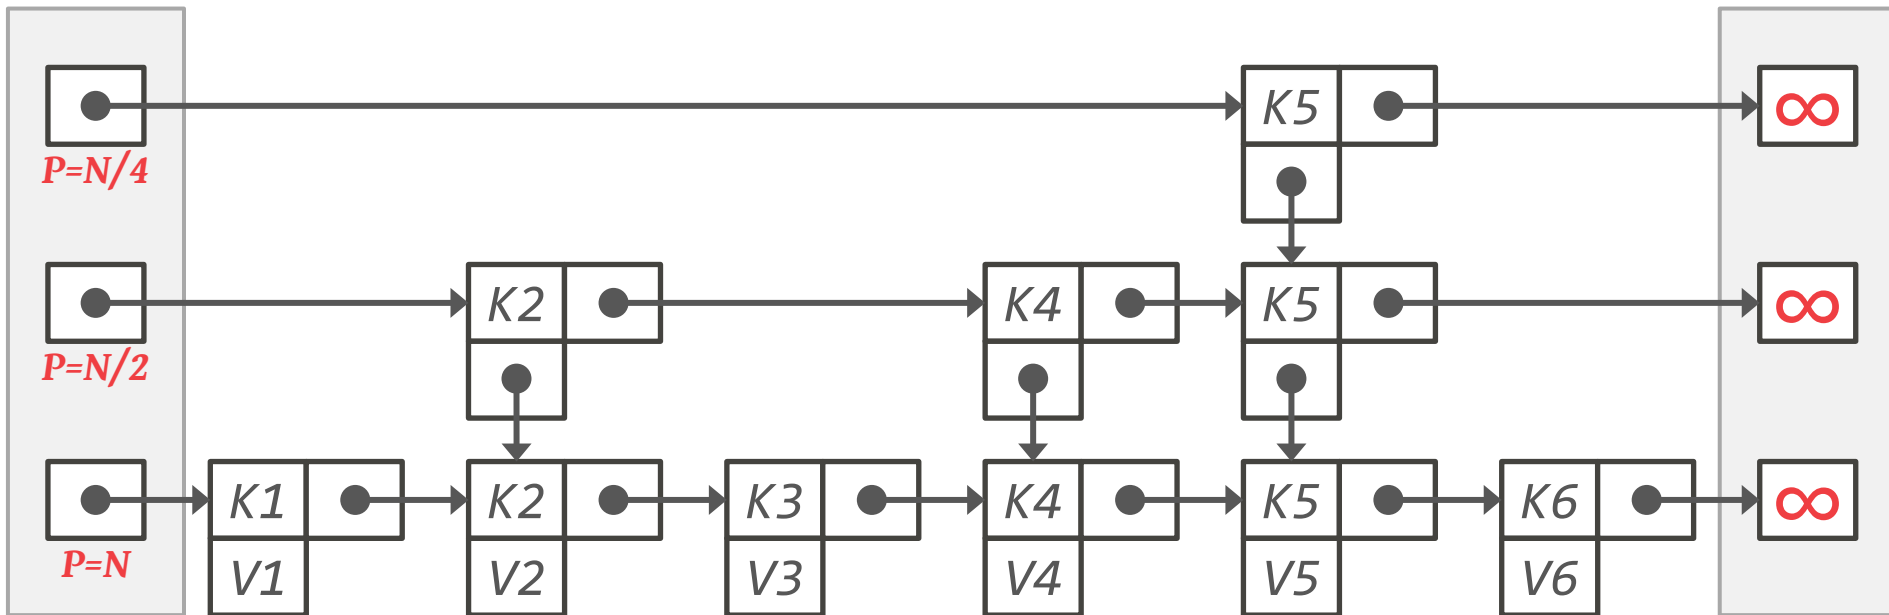


SKIP LISTS: REVERSE SEARCH

Txn #1: Find K3

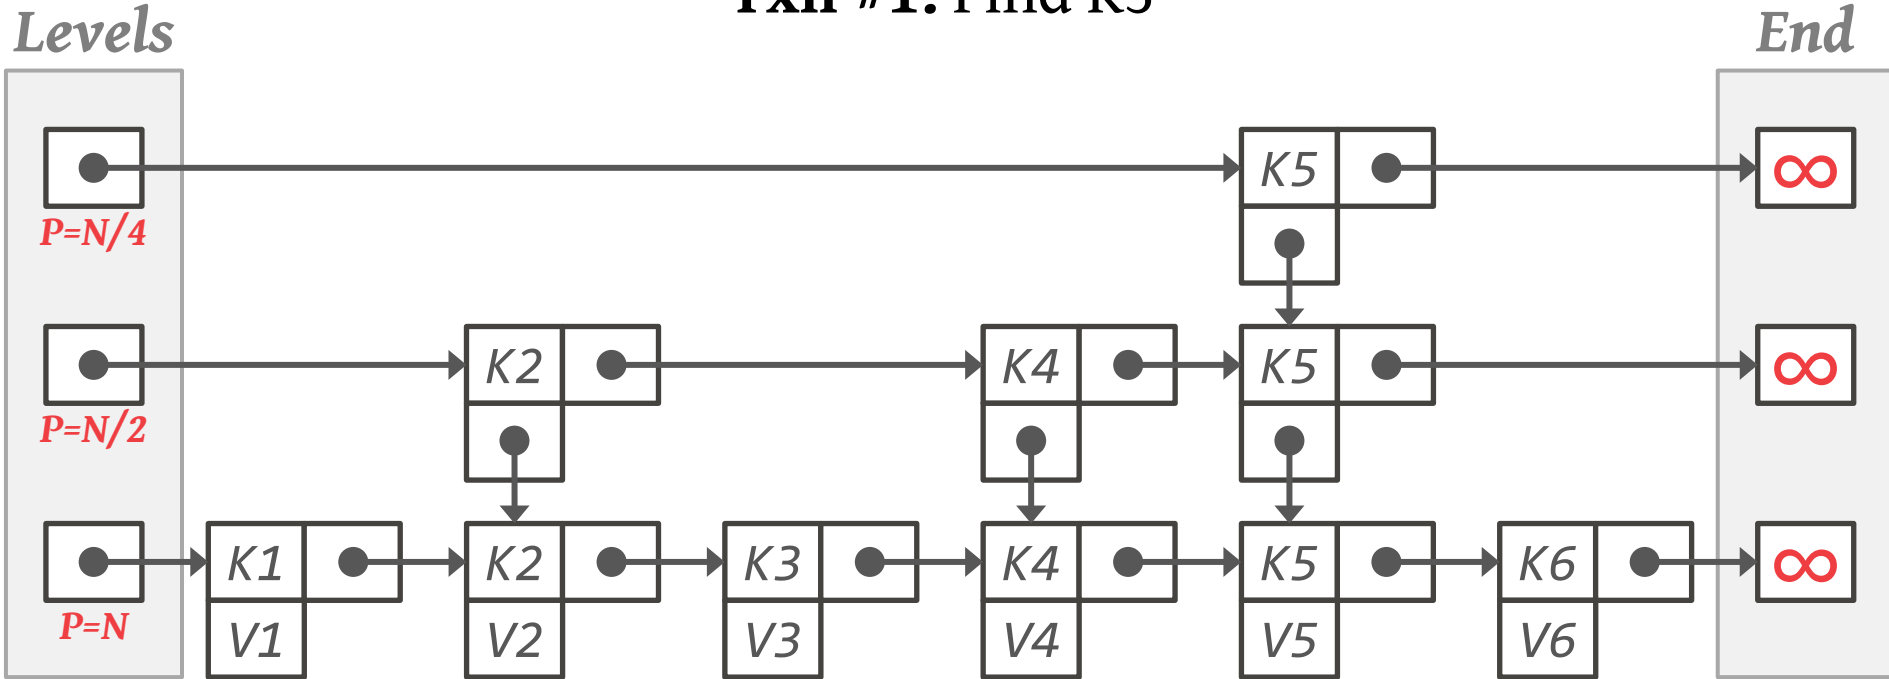
Levels

End



SKIP LISTS: REVERSE SEARCH

Txn #1: Find K3

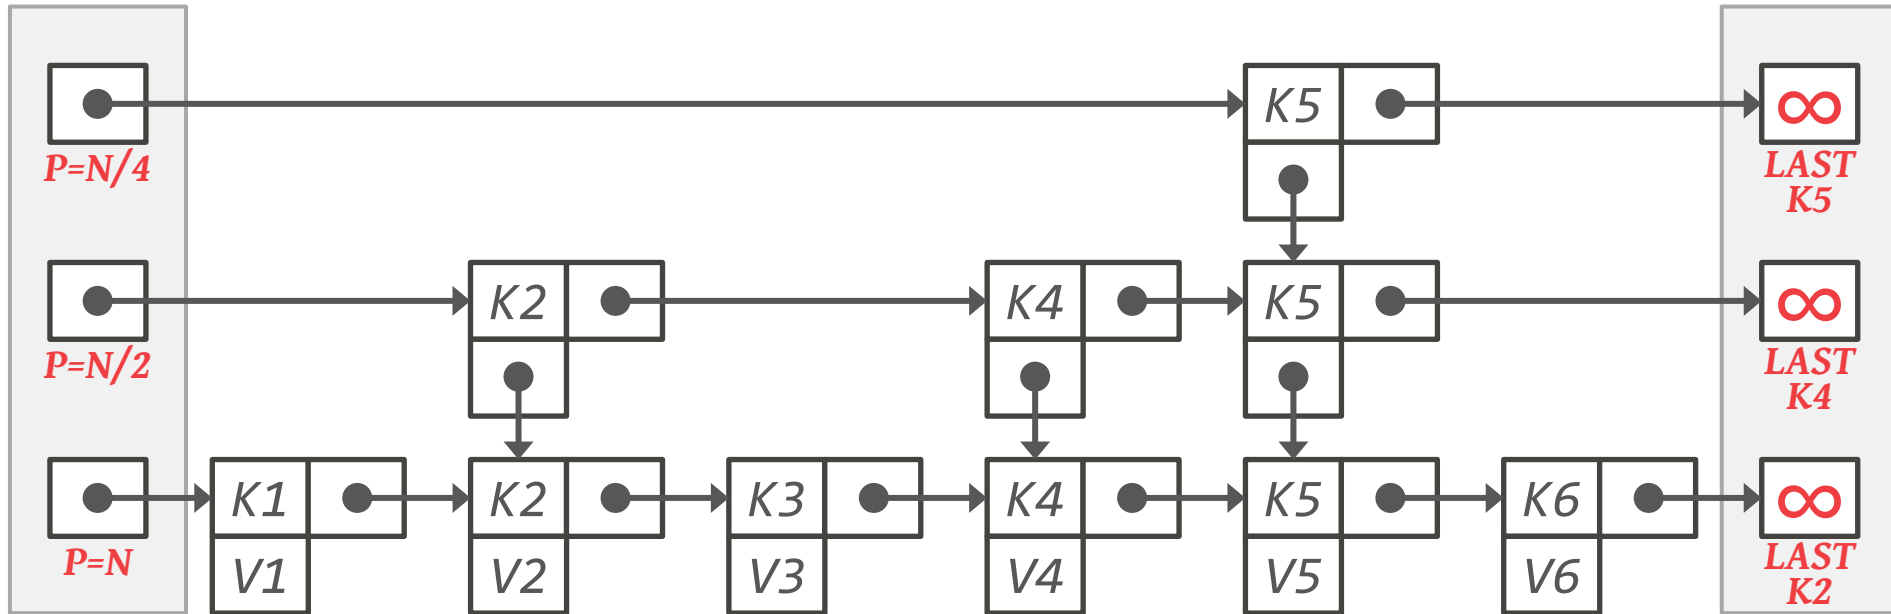


SKIP LISTS: REVERSE SEARCH

Txn #1: Find K3

Levels

End

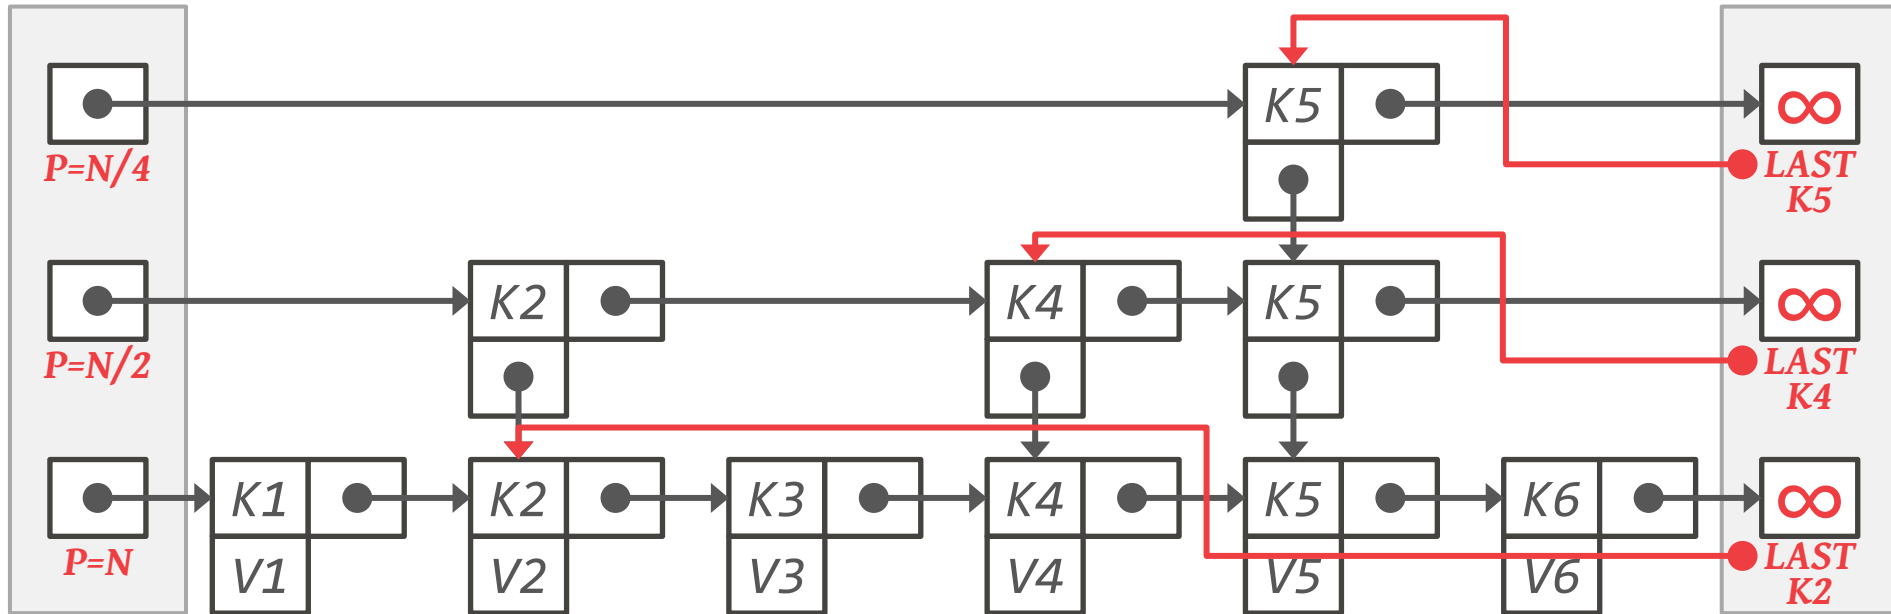


SKIP LISTS: REVERSE SEARCH

Txn #1: Find K3

Levels

End

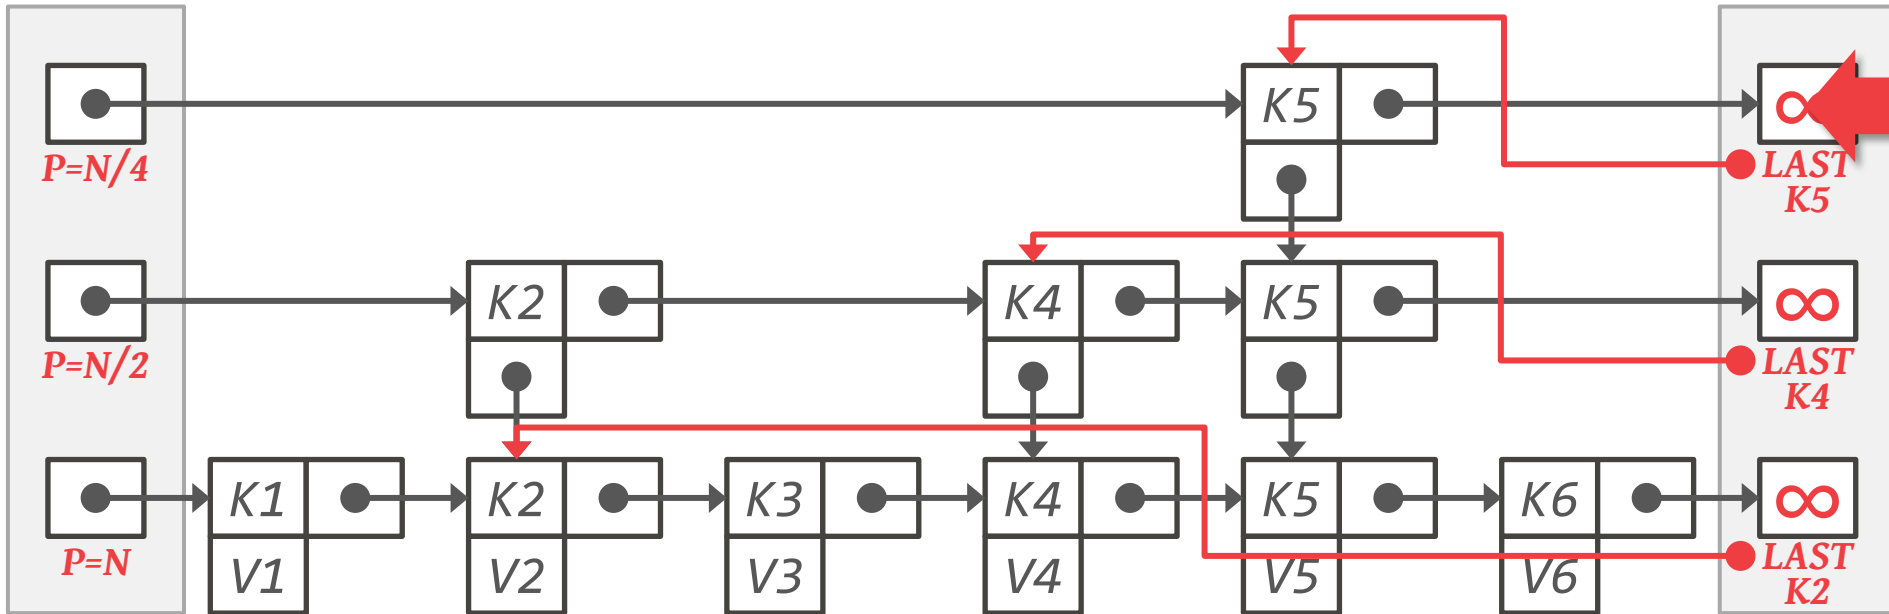


SKIP LISTS: REVERSE SEARCH

Txn #1: Find K3

Levels

End

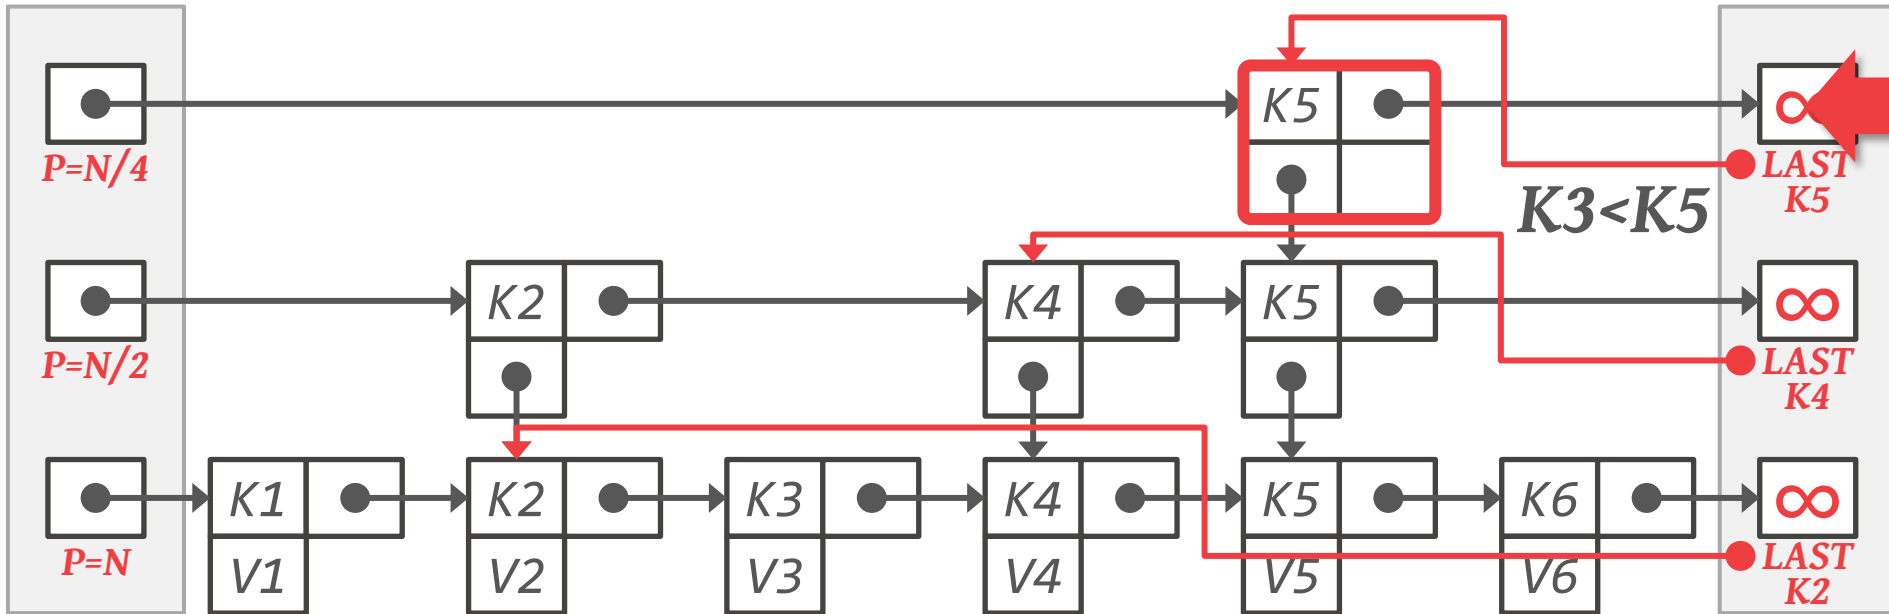


SKIP LISTS: REVERSE SEARCH

Txn #1: Find K3

Levels

End

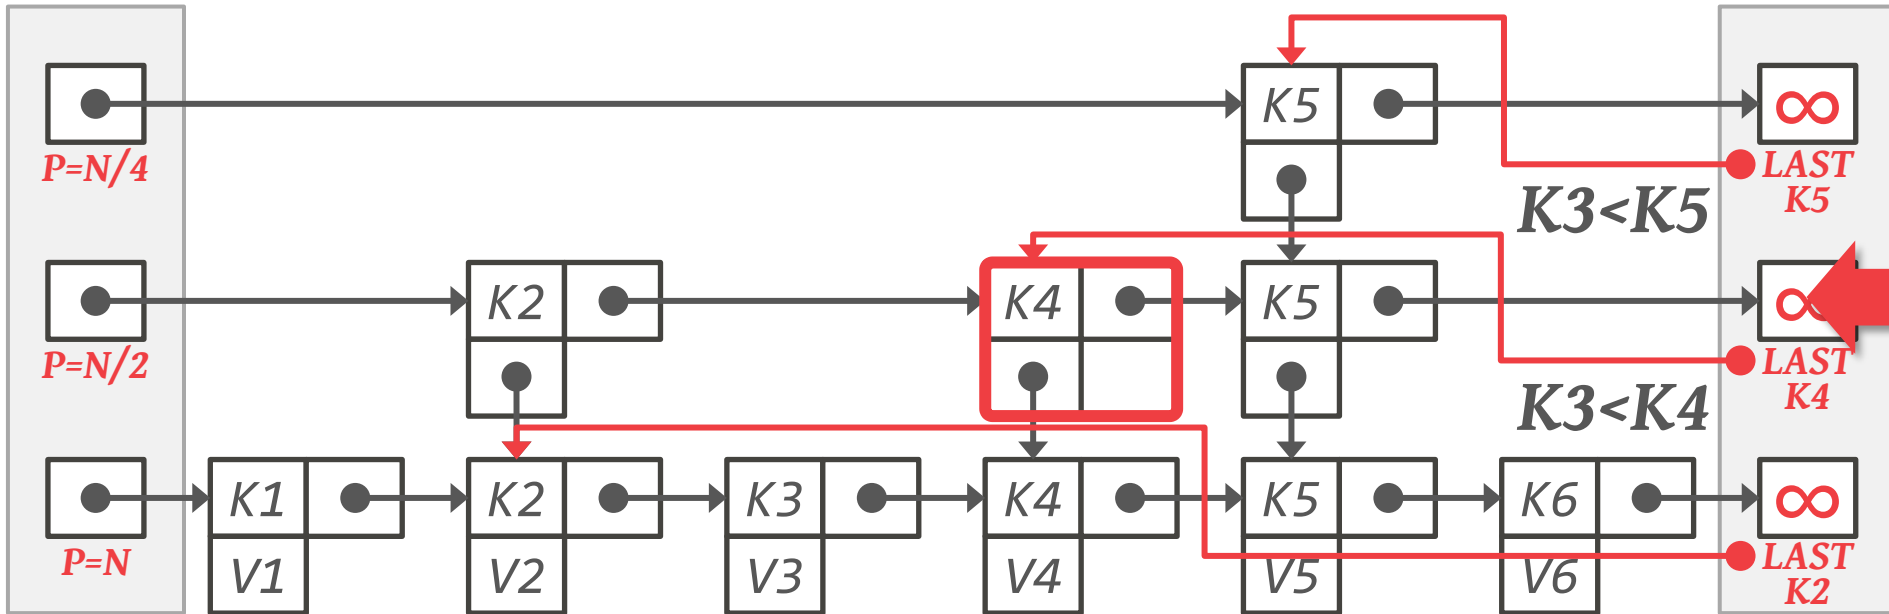


SKIP LISTS: REVERSE SEARCH

Txn #1: Find K3

Levels

End

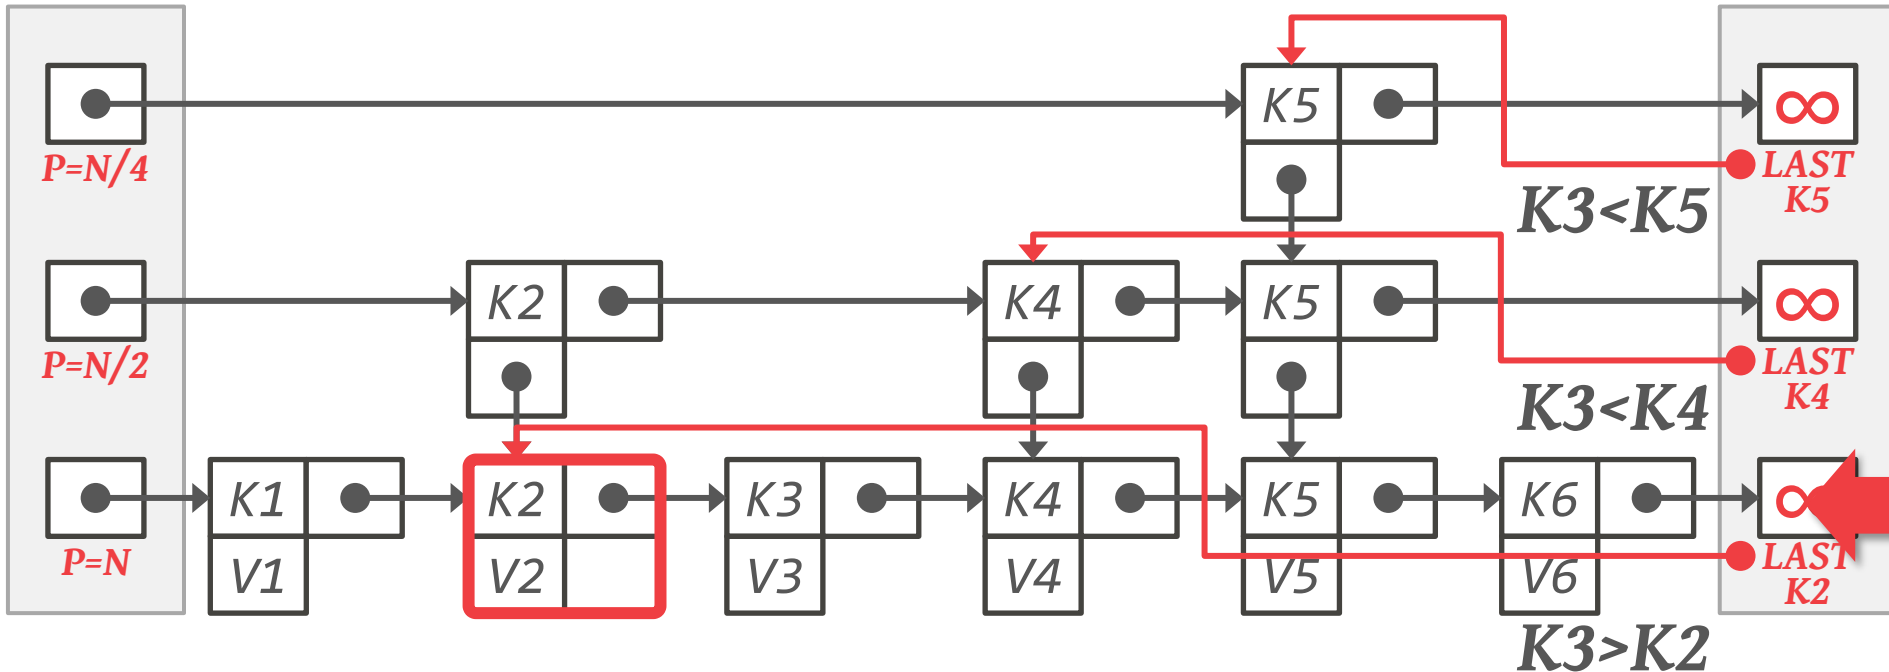


SKIP LISTS: REVERSE SEARCH

Txn #1: Find K3

Levels

End

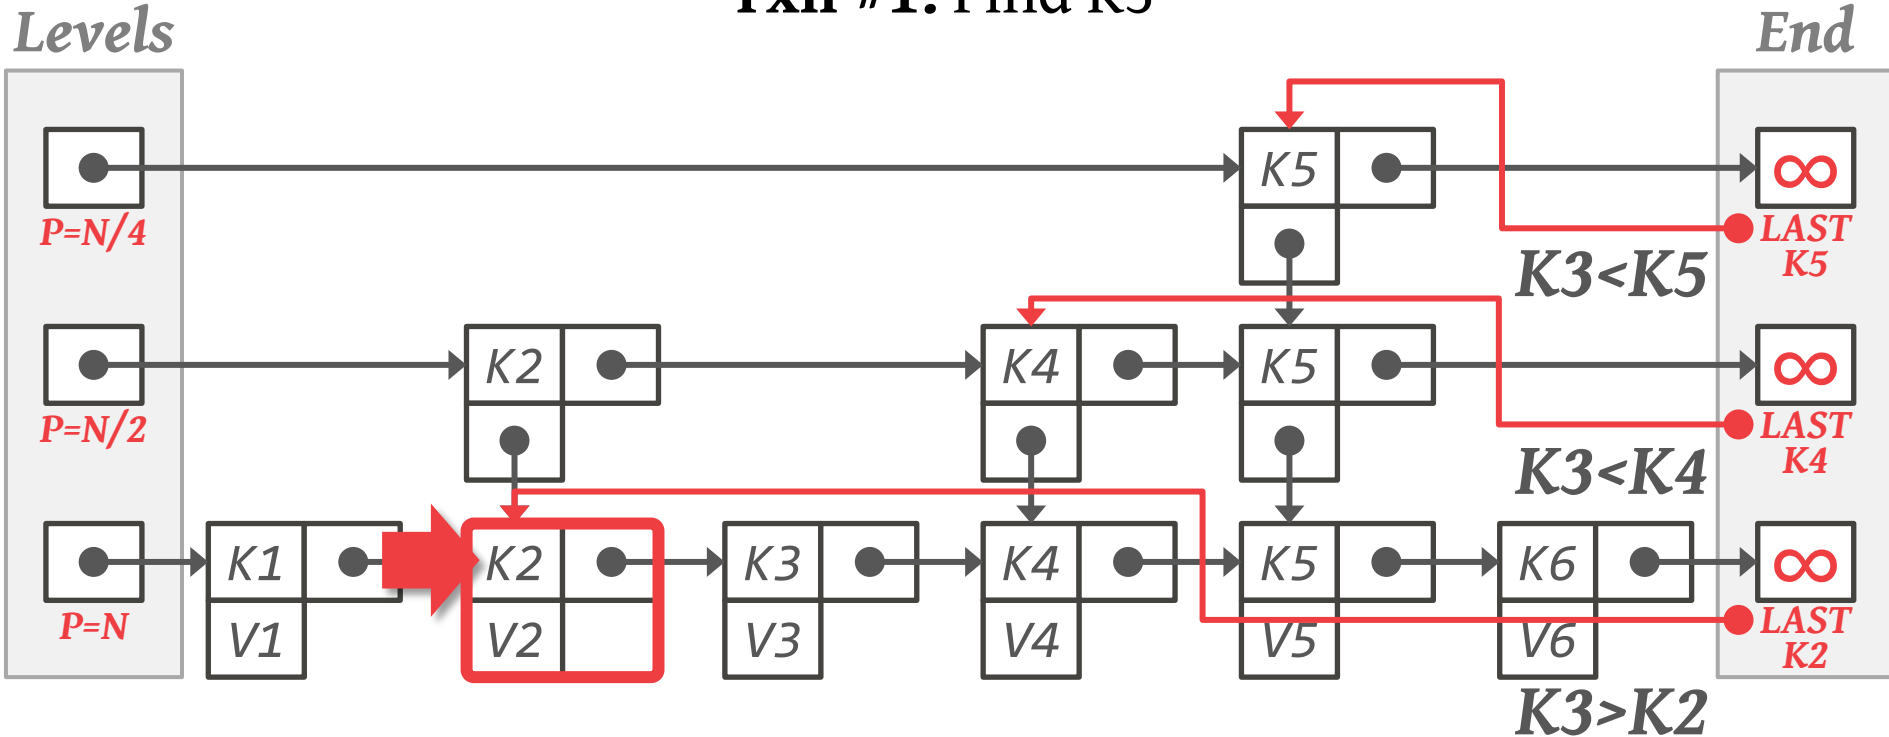


Source: [MemSQL](#)

CMU 15-721 (Spring 2016)

SKIP LISTS: REVERSE SEARCH

Txn #1: Find K3

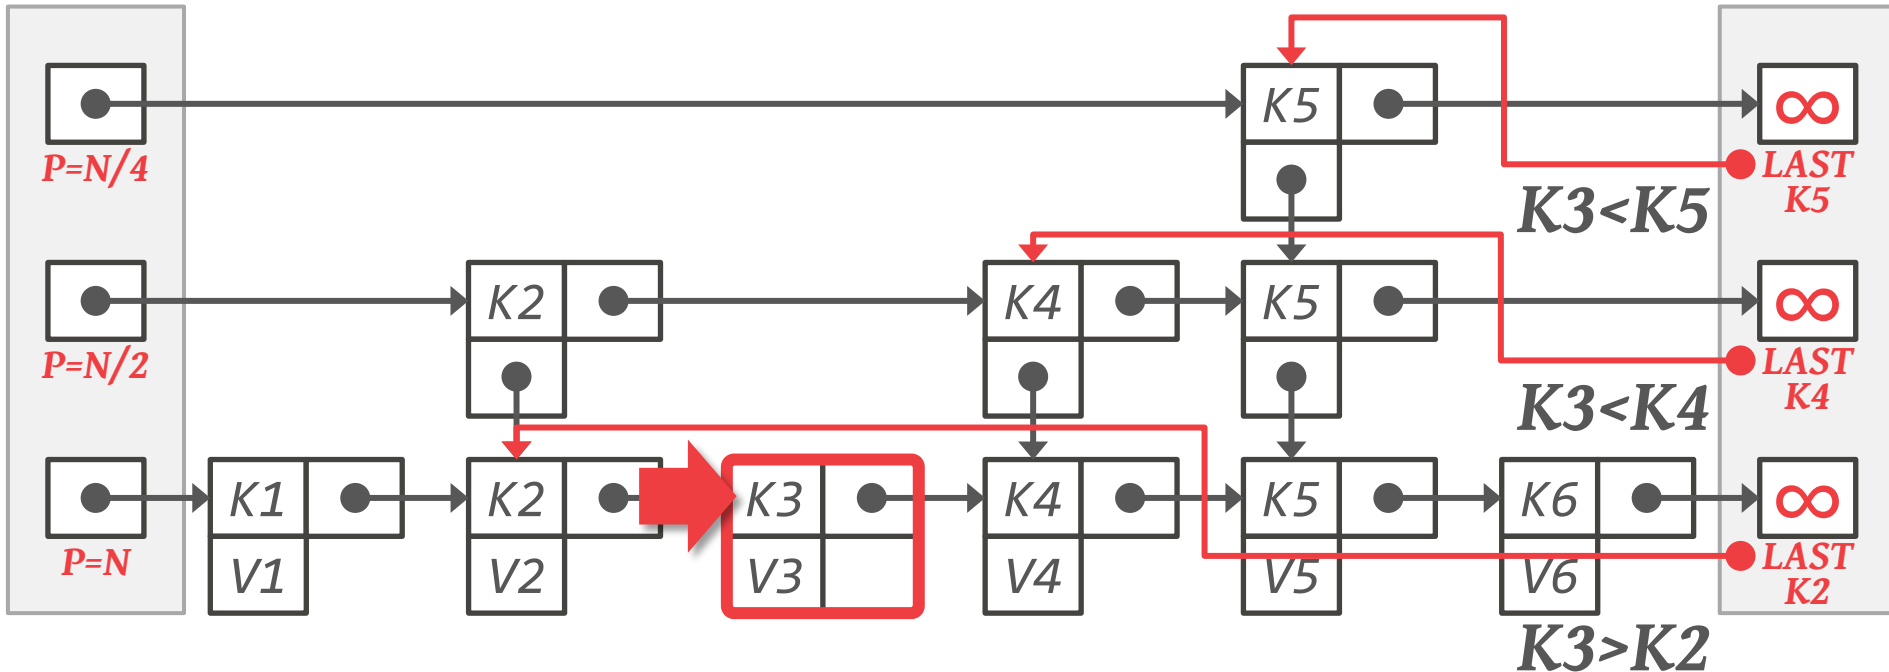


SKIP LISTS: REVERSE SEARCH

Txn #1: Find K3

Levels

End



Source: [MemSQL](#)


CMU 15-721 (Spring 2016)

WHY ARE INDEXES DIFFERENT?

The DBMS has to treat locking in indexes differently than how its concurrency control scheme manages database objects.


The physical structure can change as long as the logical contents are consistent.

PROBLEM SCENARIO #1




PROBLEM SCENARIO #1

Txn #1: Check if 25 exists




PROBLEM SCENARIO #1

Txn #1: Check if 25 exists




PROBLEM SCENARIO #1

Txn #1: Check if 25 exists




PROBLEM SCENARIO #1

Txn #1: Check if 25 exists




PROBLEM SCENARIO #1

Txn #1: Check if 25 exists




PROBLEM SCENARIO #1

Txn #1: Check if 25 exists




PROBLEM SCENARIO #1

Txn #1: Check if 25 exists




PROBLEM SCENARIO #1

Txn #1: Check if 25 exists



PROBLEM SCENARIO #1


Txn #1: Check if 25 exists



PROBLEM SCENARIO #1

Txn #1: Check if 25 exists


Txn #2: Insert 25



PROBLEM SCENARIO #1

Txn #1: Check if 25 exists


Txn #2: Insert 25



PROBLEM SCENARIO #1

Txn #1: Check if 25 exists


Txn #2: Insert 25



PROBLEM SCENARIO #1

Txn #1: Check if 25 exists

Txn #2: Insert 25




PROBLEM SCENARIO #1

Txn #1: Check if 25 exists

Txn #2: Insert 25

Txn #1: Insert 25




PROBLEM SCENARIO #1


Txn #1: Check if 25 exists

Txn #2: Insert 25

Txn #1: Insert 25




PROBLEM SCENARIO #2




PROBLEM SCENARIO #2

Txn #1: Scan [12, 23]




PROBLEM SCENARIO #2

Txn #1: Scan [12, 23]



PROBLEM SCENARIO #2


Txn #1: Scan [12, 23]



PROBLEM SCENARIO #2

Txn #1: Scan [12, 23]


Txn #2: Insert 21



PROBLEM SCENARIO #2

Txn #1: Scan [12, 23]


Txn #2: Insert 21



PROBLEM SCENARIO #2

Txn #1: Scan [12, 23]


Txn #2: Insert 21



PROBLEM SCENARIO #2

Txn #1: Scan [12, 23]

Txn #2: Insert 21




PROBLEM SCENARIO #2

Txn #1: Scan [12, 23]

Txn #2: Insert 21

Txn #1: Scan [12, 23]




PROBLEM SCENARIO #2

Txn #1: Scan [12, 23]

Txn #2: Insert 21

Txn #1: Scan [12, 23]



LOCKS VS. LATCHES

Locks

- Protects the index's logical contents from other txns.
- Held for txn duration.
- Need to be able to rollback changes.

Latches

- Protects the critical sections of the index's internal data structure from other threads.
- Held for operation duration.
- Do not need to be able to rollback changes.




LOCKS VS. LATCHES

	<i>Locks</i>	<i>Latches</i>
Separate...	User transactions	Threads
Protect...	Database Contents	In-Memory Data Structures
During...	Entire Transactions	Critical Sections
Modes...	Shared, Exclusive, Update, Intention	Read, Write
Deadlock	Detection & Resolution	Avoidance
...by...	Waits-for, Timeout, Aborts	Coding Discipline
Kept in...	Lock Manager	Protected Data Structure


INDEX LOCKS

Lock Table




INDEX LOCKS

Lock Table



INDEX LOCKS

Lock Table



LOCK-FREE INDEXES

Possibility #1: No Locks

- Txns don't acquire locks to access/modify database.
- Still have to use latches to install updates.

Possibility #2: No Latches

- Use multi-versioning inside of the index. Swap pointers using atomic updates to install updates.
- Still have to use locks to validate txns.

INDEX LOCKING

Predicate Locks

Key-Value Locks

Gap Locks

Key-Range Locks

Hierarchical Locking

PREDICATE LOCKS

Proposed locking scheme from System R.

- Shared lock on the predicate in a **WHERE** clause of a **SELECT** query.
- Exclusive lock on the predicate in a **WHERE** clause of any **UPDATE**, **INSERT** or **DELETE** query.

Never implemented in any system.



PREDICATE LOCKS

```
SELECT SUM(balance)
  FROM account
 WHERE name = 'Tupac'
```

```
INSERT INTO account
(name, balance)
VALUES ('Tupac', 100);
```



Records in Table 'account'

PREDICATE LOCKS

```
SELECT SUM(balance)
FROM account
WHERE name = 'Tupac'
```

```
INSERT INTO account
(name, balance)
VALUES ('Tupac', 100);
```



Records in Table 'account'

PREDICATE LOCKS

```
SELECT SUM(balance)
FROM account
WHERE name = 'Tupac'
```

```
INSERT INTO account
(name, balance)
VALUES ('Tupac', 100);
```



Records in Table 'account'



name='Tupac'

PREDICATE LOCKS

```
SELECT SUM(balance)
FROM account
WHERE name = 'Tupac'
```

```
INSERT INTO account
(name, balance)
VALUES ('Tupac', 100);
```



Records in Table 'account'



name='Tupac'



name='Tupac' \wedge
balance=100

KEY-VALUE LOCKS

Locks that cover a single key value.

Need “virtual keys” for non-existent values.

B+Tree Leaf Node




KEY-VALUE LOCKS

Locks that cover a single key value.

Need “virtual keys” for non-existent values.

B+Tree Leaf Node



GAP LOCKS

Each txn acquires a key-value lock on the single key that it wants to access. Then get a gap lock on the next key gap.

B+Tree Leaf Node



GAP LOCKS

Each txn acquires a key-value lock on the single key that it wants to access. Then get a gap lock on the next key gap.

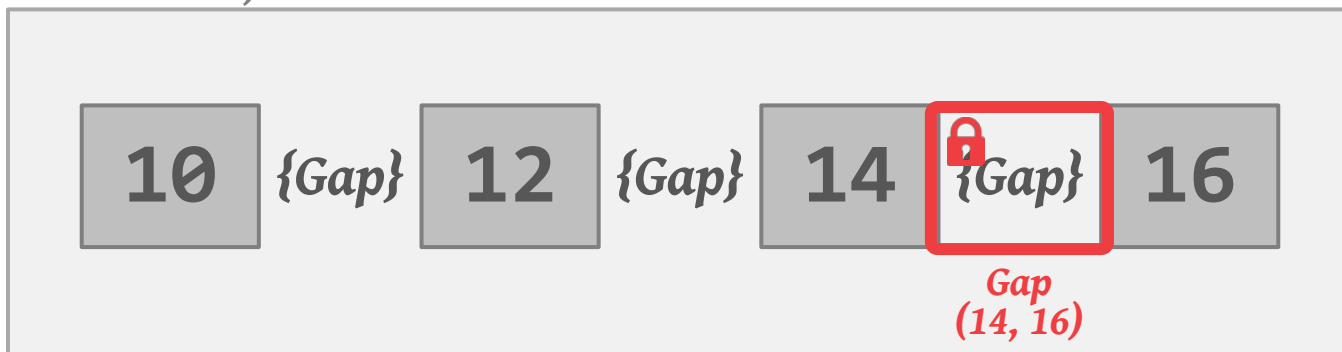
B+Tree Leaf Node



GAP LOCKS

Each txn acquires a key-value lock on the single key that it wants to access. Then get a gap lock on the next key gap.

B+Tree Leaf Node



KEY-RANGE LOCKS

A txn takes locks on ranges in the key space.


- Each range is from one key that appears in the relation, to the next that appears.
- Define lock modes so conflict table will capture commutativity of the operations available.

KEY-RANGE LOCKS

Locks that cover a key value and the gap to the next key value in a single index.

→ Need “virtual keys” for artificial values (infinity)

B+Tree Leaf Node

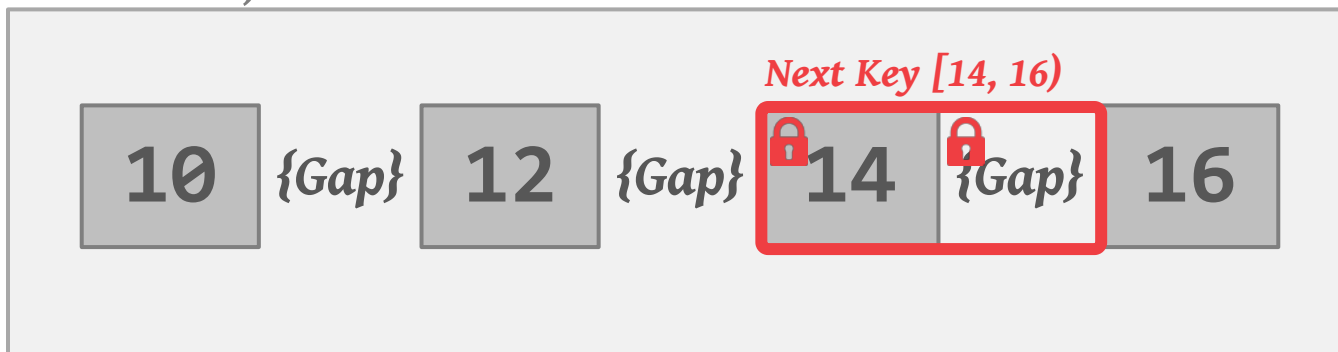


KEY-RANGE LOCKS

Locks that cover a key value and the gap to the next key value in a single index.

→ Need “virtual keys” for artificial values (infinity)

B+Tree Leaf Node

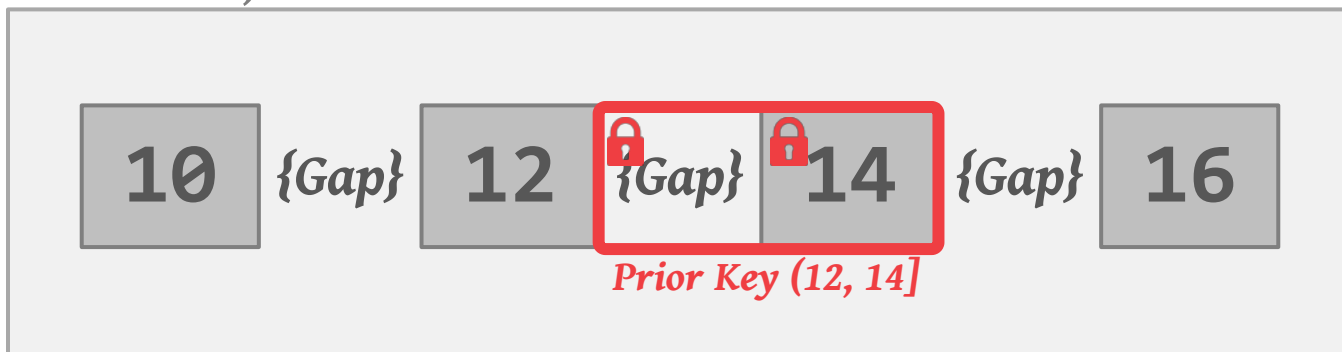


KEY-RANGE LOCKS

Locks that cover a key value and the gap to the next key value in a single index.

→ Need “virtual keys” for artificial values (infinity)

B+Tree Leaf Node

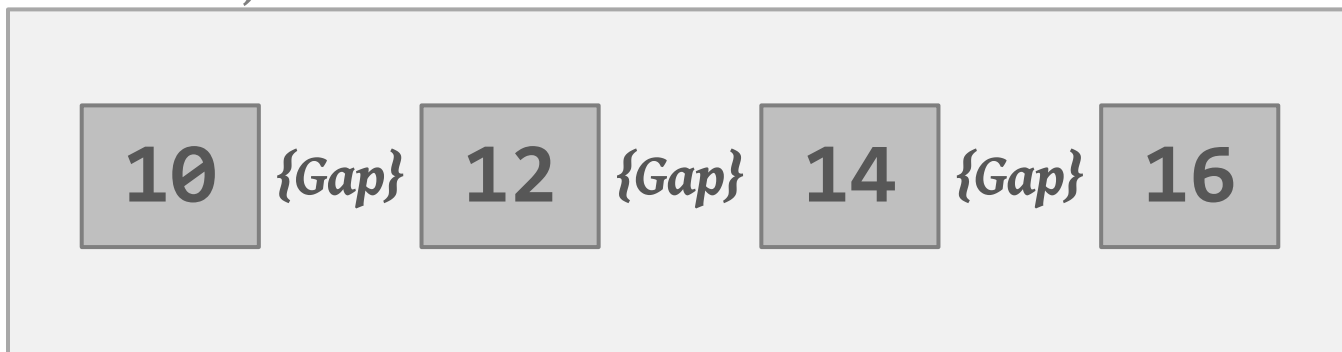


HIERARCHICAL LOCKING

Allow for a txn to hold wider key-range locks with different locking modes.

→ Reduces the number of visits to lock manager.

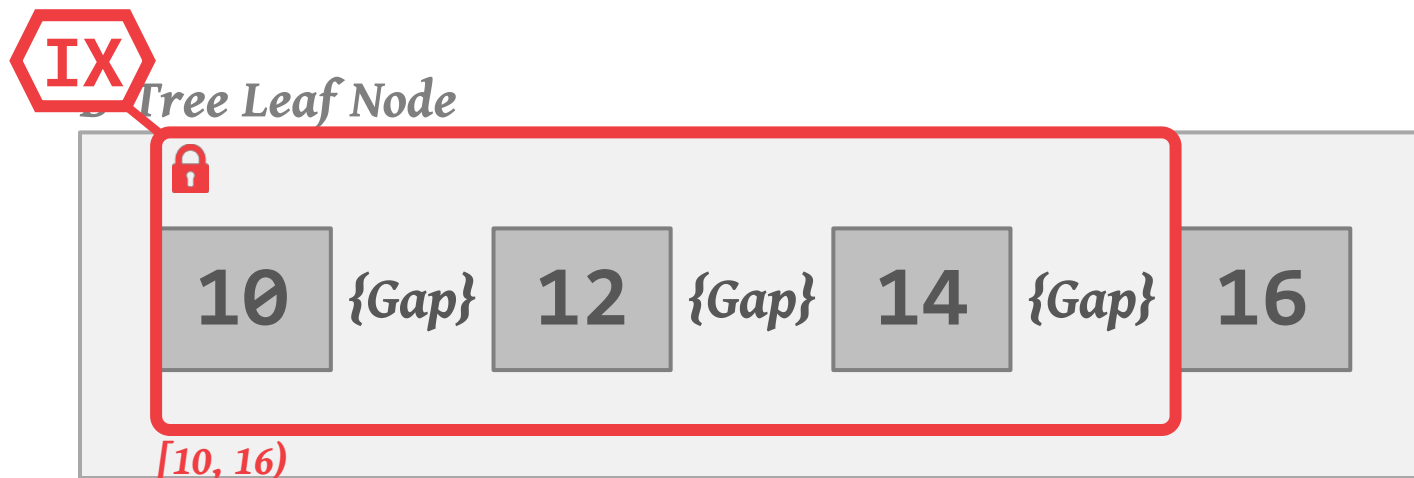
B+Tree Leaf Node



HIERARCHICAL LOCKING

Allow for a txn to hold wider key-range locks with different locking modes.


→ Reduces the number of visits to lock manager.



HIERARCHICAL LOCKING

Allow for a txn to hold wider key-range locks with different locking modes.

→ Reduces the number of visits to lock manager.



PARTING THOUGHTS

Hierarchical locking essentially provides predicate locking without complications.

- Index locking occurs only in the leaf nodes.
- Latching is to ensure consistent data structure.

Just like concurrency control schemes, research on fast indexes is hot again.

PRISON TATTOOS

Some of you are going to end up in prison.
→ This is just the nature of the database game.

Part of surviving prison is being able to
navigate and avoid the various factions.

TEAR DROP



THREE DOTS



FIVE DOTS



MARA SALVATRUCHA GANG (MS13)



ARYAN BROTHERHOOD



NEXT CLASS

Bw-Tree (Hekaton)

Concurrent Skip Lists (MemSQL)

ART Index (HyPer)